

Table of Contents

1 MIDS - w261 Machine Learning At Scale

1.1 Assignment - HW3

2 Table of Contents

3 1 Instructions

3.0.1 IMPORTANT

3.0.2 === INSTRUCTIONS for SUBMISSIONS ===

3.1 HW3.0.

3.2 HW3.1 consumer complaints dataset: Use Counters to do EDA (exploratory data analysis and to monitor progress)

3.2.1 HW 3.2 Analyze the performance of your Mappers, Combiners and Reducers using Counters

3.2.2 3.2.A SOLUTION

3.2.2.1 3.2.A EXPLANATION

3.2.3 3.2.B SOLUTION

3.2.4 3.2.C SOLUTION

3.2.5 3.2.1

3.2.5.1 START STUDENT CODE HW321 (INSERT CELLS BELOW AS NEEDED)

3.2.5.2 END STUDENT CODE HW321

3.3 HW3.3. Shopping Cart Analysis

3.3.0.1 START STUDENT CODE HW33 (INSERT CELLS BELOW AS NEEDED)

3.3.0.2 END STUDENT CODE HW33

3.4 HW3.3.1 OPTIONAL

3.4.0.1 START STUDENT CODE HW331 (INSERT CELLS BELOW AS NEEDED)

3.4.0.2 END STUDENT CODE HW331

3.5 HW3.4. (Computationally prohibitive but then again Hadoop can handle this) Pairs

3.5.0.1 START STUDENT CODE HW34 (INSERT CELLS BELOW AS NEEDED)

3.5.0.2 END STUDENT CODE HW34

3.6 HW3.5: Stripes

3.6.0.1 START STUDENT CODE HW35 (INSERT CELLS BELOW AS NEEDED)

3.6.0.2 END STUDENT CODE HW35

4 OPTIONAL

4.1 HW3.6 Computing Relative Frequencies on 100K WikiPedia pages (93Meg)

4.2 HW3.7 Apriori Algorithm

4.3 HW3.8. Shopping Cart Analysis

4.4 HW3.8.1

5 END OF HOMEWORK

```
In [1]: %%javascript
/*****
*****
Known Mathjax Issue with Chrome - a rounding issue adds a border to the right o
f mathjax markup
https://github.com/mathjax/MathJax/issues/1300
A quick hack to fix this based on stackoverflow discussions:
http://stackoverflow.com/questions/34277967/chrome-rendering-mathjax-equations-
with-a-trailing-vertical-line
*****/

$('.math>span').css("border-left-color","transparent")
```

```
In [2]: %reload_ext autoreload
%autoreload 2
```

MIDS - w261 Machine Learning At Scale

Course Lead: Dr James G. Shanahan (**email** Jimi via James.Shanahan AT gmail.com)

Assignment - HW3

Name: Nilesh Bhoyat **Class:** MIDS w261 Summer Group 2

Email: *Your UC Berkeley Email Goes Here*@iSchool.Berkeley.edu

StudentId 26302327 **End of StudentId**

Week: 3

NOTE: please replace 1234567 with your student id above

Due Time: HW is due the Tuesday of the following week by 8AM (West coast time). I.e., Tuesday, Jan 31, 2017 in the case of this homework.

Table of Contents

1. [HW Intructions](#)
2. [HW Problems](#)
 - 3.1. [HW3.1](#)
 - 3.2. [HW3.2](#)
 - 3.3 [HW3.3](#)
 - 3.4 [HW3.4](#)
 - 3.5 [HW3.5](#)
 - 3.6 [HW3.6](#)
 - 3.7 [HW3.7](#)
 - 3.8 [HW3.8](#)

1 Instructions

[Back to Table of Contents](#)

MIDS UC Berkeley, Machine Learning at Scale DATSCIW261 ASSIGNMENT #3

Version 2017-26-1

IMPORTANT

This homework can be completed locally on your computer

=== INSTRUCTIONS for SUBMISSIONS ===

Follow the instructions for submissions carefully.

NEW: Going forward, each student will have a `HW-<user>` repository for all assignments.

Click this link to enable you to create a github repo within the MIDS261 Classroom:

<https://classroom.github.com/assignment-invitations/3b1d6c8e58351209f9dd865537111ff8>
[\(https://classroom.github.com/assignment-invitations/3b1d6c8e58351209f9dd865537111ff8\)](https://classroom.github.com/assignment-invitations/3b1d6c8e58351209f9dd865537111ff8)

and follow the instructions to create a HW repo.

Push the following to your HW github repo into the master branch:

- Your local HW3 directory. Your repo file structure should look like this:

```
HW-<user>
  --HW3
    |__MIDS-W261-HW-03-<Student_id>.ipnb
    |__MIDS-W261-HW-03-<Student_id>.pdf
    |__some other hw3 file
  --HW4
    |__MIDS-W261-HW-04-<Student_id>.ipnb
    |__MIDS-W261-HW-04-<Student_id>.pdf
    |__some other hw4 file
  etc..
```

HW3.0.

1. How do you merge two sorted lists/arrays of records of the form [key, value]?
2. Where is this used in Hadoop MapReduce? [Hint within the shuffle]
3. What is a combiner function in the context of Hadoop?
4. Give an example where it can be used and justify why it should be used in the context of this problem.
5. What is the Hadoop shuffle?

How do you merge two sorted lists/arrays of records of the form [key, value]?

Merge sort is divide and conquer algorithm. It is a very efficient sort algorithm. The algorithm gets its name from the fact that it divides the collection in half, recursively sorts each half, and then merges the two sorted halves back together. Each half of the collection is repeatedly halved until there is only one object in the half, at which point it is sorted by definition. As each sorted half is merged, the algorithm compares the objects to determine where to place each sub set.

for records in format [key,value], operations happen on keys.
Pseudo Code

Union of keys from both dictionaries.
Loop for each key in union:
 if key is in dict1 and not in dict 2 then
 add key records in result dict
 if key is in dict2 and not in dict 1 then
 add key record in result dict
 if key is present in both dicts then
 union of result with same key

```
In [3]: from collections import defaultdict
dict1 = {'bookA': 1, 'bookB': 2, 'bookC': 3}
dict2 = {'bookC': 2, 'bookD': 4, 'bookE': 5}

def union_collections(d1, d2):
    union = {}

    for key in set(d1.keys()).union(d2.keys()):
        if key in d1 and key not in d2:
            union[key] = d1[key]

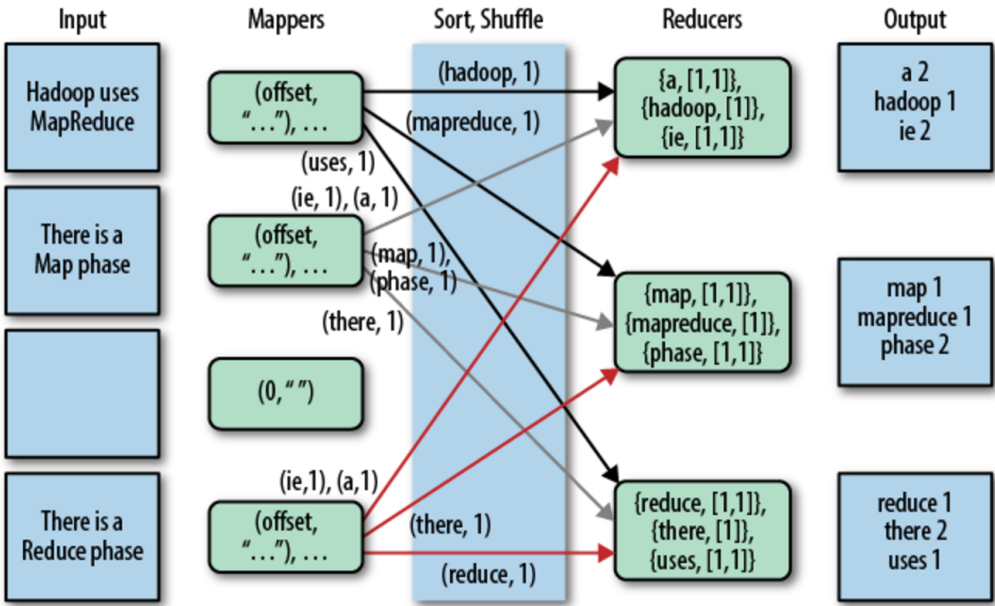
        if key in d2 and key not in d1:
            union[key] = d2[key]

        if key in d1 and key in d2:
            union[key] = (d1[key], d2[key])

    return union
union_collections(dict1, dict2)
```

```
Out[3]: {'bookA': 1, 'bookB': 2, 'bookC': (3, 2), 'bookD': 4, 'bookE': 5}
```

Where is this used in Hadoop MapReduce? [Hint within the shuffle]



Answer:
As shown in figure, key-value pairs are merge-sort is used in shuffle-sort phase. Hadoop sorts the key-value pairs by key and it "shuffles" all pairs with the same key to the same Reducer. There are several possible techniques that can be used to decide which reducer gets which range of keys.

What is a combiner function in the context of Hadoop?

Mappers produce a lot of intermediate data that must be sent over the network to be shuffled, sorted, and reduced. Because networking is a physical resource, large amounts of transmitted data can lead to job delays and memory bottlenecks (e.g., there is too much data for the reducer to hold into memory). Combiners are the primary mechanism to solve this problem, and are essentially intermediate reducers that are associated with the mapper output. Combiners reduce network traffic by performing a mapper-local reduction of the data before forwarding it on to the appropriate reducer.

Give an example where it can be used and justify why it should be used in the context of this problem.

Answer

Consider an example of wordcount for large corpora. When multiple mappers are producing word count as below

mapper 1 output

(a,10) (the,10) (tent,1) (the,20)(a,20)

Mapper 2 output

(a,30) (the,10)

Intended Sum Reduce is

(a,60) (the,40)(tent,1)

Each mapper is emitting extra work for the reducer, namely in the duplication of the different keys coming from each mapper. Combiner can reduce such duplicate key by aggregating records before and thereby reduce network traffic. This will also reduce overall shuffle phase time.

What is the Hadoop shuffle?

Answer

The process by which the Hadoop performs the sort-and transfers the map outputs to the reducers as inputs.

If reducer has to take all output of mapper in as-is format i.e. (word,1) then it would be a very slow process.

Shuffle does this

Shuffle is all of this!:

1. partition, sort, combine - Partitions records and does partial sort for each partition.

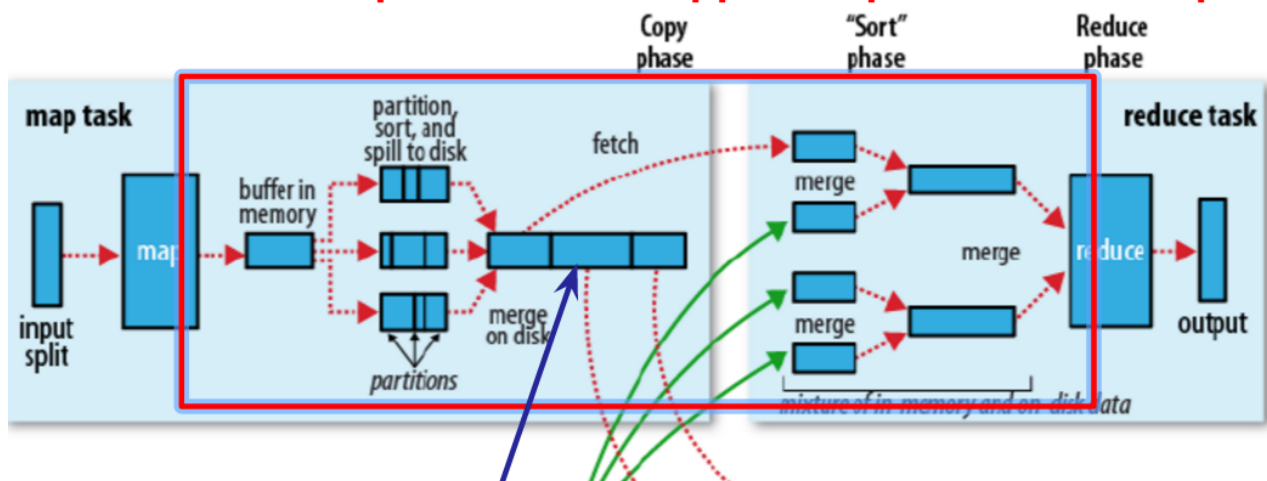
2. mergesort

3. Send to reducer

4. Merge sort

5. Stream to reducer

Shuffle: all steps between mapper output to reduce input



HW3.1 consumer complaints dataset: Use Counters to do EDA (exploratory data analysis and to monitor progress)

Counters are lightweight objects in Hadoop that allow you to keep track of system progress in both the map and reduce stages of processing. By default, Hadoop defines a number of standard counters in "groups"; these show up in the jobtracker webapp, giving you information such as "Map input records", "Map output records", etc.

While processing information/data using MapReduce job, it is a challenge to monitor the progress of parallel threads running across nodes of distributed clusters. Moreover, it is also complicated to distinguish between the data that has been processed and the data which is yet to be processed. The MapReduce Framework offers a provision of user-defined Counters, which can be effectively utilized to monitor the progress of data across nodes of distributed clusters.

Use the Consumer Complaints Dataset provide here to complete this question:

https://www.dropbox.com/s/vbalm3yva2rr86m/Consumer_Complaints.csv?dl=0

The consumer complaints dataset consists of diverse consumer complaints, which have been reported across the United States regarding various types of loans. The dataset consists of records of the form:

Complaint ID,Product,Sub-product,Issue,Sub-issue,State,ZIP code,Submitted via,Date received,Date sent to company,Company,Company response,Timely response?,Consumer disputed?

Here's is the first few lines of the of the Consumer Complaints Dataset:

```
Complaint ID,Product,Sub-product,Issue,Sub-issue,State,ZIP code,Submitted via,Date received,Date sent to company,Company,Company response,Timely response?,Consumer disputed?
1114245,Debt collection,Medical,Disclosure verification of debt,Not given enough in fo to verify debt,FL,32219,Web,11/13/2014,11/13/2014,"Choice Recovery, Inc.",Closed with explanation,Yes,
1114488,Debt collection,Medical,Disclosure verification of debt,Right to dispute no tice not received,TX,75006,Web,11/13/2014,11/13/2014,"Expert Global Solutions, Inc.",In progress,Yes,
1114255,Bank account or service,Checking account,Deposits and withdrawals,,NY,11102,Web,11/13/2014,11/13/2014,"FNIS (Fidelity National Information Services, Inc.)",In progress,Yes,
1115106,Debt collection,"Other (phone, health club, etc.)",Communication tactics,Fr equent or repeated calls,GA,31721,Web,11/13/2014,11/13/2014,"Expert Global Solution s, Inc.",In progress,Yes,
```

User-defined Counters

Now, let's use Hadoop Counters to identify the number of complaints pertaining to debt collection, mortgage and other categories (all other categories get lumped into this one) in the consumer complaints dataset. Basically produce the distribution of the Product column in this dataset using counters (limited to 3 counters here).

Hadoop offers Job Tracker, an UI tool to determine the status and statistics of all jobs. Using the job tracker UI, developers can view the Counters that have been created. Screenshot your job tracker UI as your job completes and include it here. Make sure that your user defined counters are visible.

```
In [4]: # Put the data into HDFS
!wget 'https://www.dropbox.com/s/vbalm3yva2rr86m/Consumer_Complaints.csv'

--2017-05-30 05:53:50-- https://www.dropbox.com/s/vbalm3yva2rr86m/Consumer_Co
mplaints.csv
Resolving www.dropbox.com... 162.125.4.1
Connecting to www.dropbox.com|162.125.4.1|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://dl.dropboxusercontent.com/content_link/bySujatrsT8oJft8X7tWO
kDd81kv2gy0533Cjc3drnm5QzbM5ET3cDIz62ValFT/file [following]
--2017-05-30 05:53:57-- https://dl.dropboxusercontent.com/content_link/bySuja
trst8oJft8X7tWOkDd81kv2gy0533Cjc3drnm5QzbM5ET3cDIz62ValFT/file
Resolving dl.dropboxusercontent.com... 162.125.4.6
Connecting to dl.dropboxusercontent.com|162.125.4.6|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 50906486 (49M) [text/csv]
Saving to: `Consumer_Complaints.csv'

100%[=====>] 50,906,486 1.23M/s in 92s

2017-05-30 05:55:32 (541 KB/s) - `Consumer_Complaints.csv' saved [50906486/509
06486]
```

```
In [5]: # Create HDFS directories
!hdfs dfs -mkdir -p /user/nibhoyar
!hdfs dfs -put Consumer_Complaints.csv /user/nibhoyar
!hdfs dfs -rm Consumer_Complaints.csv
!hdfs dfs -copyFromLocal Consumer_Complaints.csv
!hdfs dfs -rm -r hw3.1-output

put: `/user/nibhoyar/Consumer_Complaints.csv': File exists
Deleted Consumer_Complaints.csv
Deleted hw3.1-output
```

```
In [6]: %%writefile complaintCountsMapper.py
#!/usr/bin/env python
# START STUDENT CODE HW31MAPPER
import sys
separator = ','
for line in (sys.stdin):
    fields = line.split(separator)
    if 'Complaint ID' != fields[0] :

        # we have a real record, so do some mapping
        counter_name = None
        if (fields[1].lower() == 'debt collection' or \
            fields[1].lower() == 'mortgage'):
            counter_name = fields[1].strip().lower()
        else:
            counter_name = 'other'
        # update the counter
        sys.stderr.write("reporter:counter:Category Counters,{0},1\n".forma
t(counter_name))

# END STUDENT CODE HW31MAPPER

Overwriting complaintCountsMapper.py
```

```
In [7]: !chomd a+x complaintCountsMapper.py

/bin/sh: chomd: command not found
```



```
In [8]: %%writefile complaintCountsReducer.py
        #!/usr/bin/env python
        # START STUDENT CODE HW31REDUCER

        # END STUDENT CODE HW31REDUCER

Overwriting complaintCountsReducer.py
```

```
In [9]: # Hadoop command
# START STUDENT CODE HW31HADOOP
!hdfs dfs -rm -r hw3.1-output
!hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
    -D mapred.reduce.tasks=0 \
    -files complaintCountsMapper.py \
    -mapper complaintCountsMapper.py \
    -reducer org.apache.hadoop.mapred.lib.IdentityReducer \
    -input Consumer_Complaints.csv \
    -output hw3.1-output
# END STUDENT CODE HW31HADOOP
```

```

rm: `hw3.1-output': No such file or directory
17/05/30 05:56:00 INFO Configuration.deprecation: mapred.reduce.tasks is depre
cated. Instead, use mapreduce.job.reduces
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob1644182036203141886.jar tmpDir=null
17/05/30 05:56:02 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:56:03 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:56:06 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 05:56:07 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 05:56:08 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0086
17/05/30 05:56:09 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0086
17/05/30 05:56:09 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0086/
17/05/30 05:56:09 INFO mapreduce.Job: Running job: job_1496033164706_0086
17/05/30 05:56:25 INFO mapreduce.Job: Job job_1496033164706_0086 running in ub
er mode : false
17/05/30 05:56:25 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 05:56:36 INFO mapreduce.Job: map 50% reduce 0%
17/05/30 05:56:37 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 05:56:37 INFO mapreduce.Job: Job job_1496033164706_0086 completed suc
cessfully
17/05/30 05:56:38 INFO mapreduce.Job: Counters: 33
    File System Counters
        FILE: Number of bytes read=0
        FILE: Number of bytes written=232896
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50910816
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=10
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=18446
        Total time spent by all reduces in occupied slots (ms)=0
        Total time spent by all map tasks (ms)=18446
        Total vcore-seconds taken by all map tasks=18446
        Total megabyte-seconds taken by all map tasks=18888704
    Map-Reduce Framework
        Map input records=312913
        Map output records=0
        Input split bytes=234
        Spilled Records=0
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=139
        CPU time spent (ms)=4320
        Physical memory (bytes) snapshot=392908800
        Virtual memory (bytes) snapshot=2711154688
        Total committed heap usage (bytes)=367525888
    Category Counters
        debt collection=44372
        mortgage=125752
        other=142788
    File Input Format Counters
        Bytes Read=50910582
    File Output Format Counters
        Bytes Written=0
17/05/30 05:56:38 INFO streaming.StreamJob: Output directory: hw3.1-output

```

output

	Name	Map	Reduce	Total
Map-Reduce Framework	CPU time spent (ms)	4400	0	4400
	Failed Shuffles	0	0	0
	GC time elapsed (ms)	180	0	180
	Input split bytes	234	0	234
	Map input records	312913	0	312913
	Map output records	0	0	0
	Merged Map outputs	0	0	0
	Physical memory (bytes) snapshot	397471744	0	397471744
	Spilled Records	0	0	0
	Total committed heap usage (bytes)	462946304	0	462946304
	Virtual memory (bytes) snapshot	2720112640	0	2720112640
	Name	Map	Reduce	Total
Category Counters	debt collection	44372	0	44372
	mortgage	125752	0	125752
	other	142788	0	142788

HW 3.2 Analyze the performance of your Mappers, Combiners and Reducers using Counters

For this brief study the Input file will be one record (the next line only):

```
foo foo quux labs foo bar quux
```

3.2.A

Perform a word count analysis of this single record dataset using a Mapper and Reducer based WordCount (i.e., no combiners are used here) using user defined Counters to count up how many times the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing this word count job. The answer should be 1 and 4 respectively. Please explain.

3.2.B

Please use multiple mappers and reducers for these jobs (at least 2 mappers and 2 reducers). Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper and Reducer based WordCount (i.e., no combiners used anywhere) using user defined Counters to count up how many time the mapper and reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job.

3.2.C

Perform a word count analysis of the Issue column of the Consumer Complaints Dataset using a Mapper, Reducer, and standalone combiner (i.e., not an in-memory combiner) based WordCount using user defined Counters to count up how many time the mapper, combiner, reducer are called. What is the value of your user defined Mapper Counter, and Reducer Counter after completing your word count job.

Using a single reducer:

- What are the top 50 most frequent terms in your word count analysis?
- Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order.
- Present bottom 10 tokens (least frequent items).

NOTE: You can use: `WORD_RE = re.compile(r"[\w']+")` to tokenize.

3.2.A SOLUTION

```
In [10]: %%writefile mapper3.2.A.py
#!/usr/bin/env python
# START STUDENT CODE HW32AMAPPER
import sys
import re

sys.stderr.write("reporter:counter:Mapper Counters,Calls,1\n")
WORD_RE = re.compile(r"[\w']+")
for line in sys.stdin:
    for word in [s.lower() for s in WORD_RE.findall(line)]:
        print '%s\t%s' % (word, 1)

# END STUDENT CODE HW32AMAPPER
```

Overwriting mapper3.2.A.py

```
In [11]: %%writefile reducer3.2.A.py
#!/usr/bin/env python
# START STUDENT CODE HW32AREducer
import sys

cur_key = None
cur_count = 0
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")
for line in sys.stdin:
    key, value = line.split()
    if key == cur_key:
        cur_count += int(value)
    else:
        if cur_key:
            print '%s\t%s' % (cur_key, cur_count)
        cur_key = key
        cur_count = int(value)

print '%s\t%s' % (cur_key, cur_count)

# END STUDENT CODE HW32AREducer
```

Overwriting reducer3.2.A.py

```
In [12]: !echo "foo foo quux labs foo bar quux"|python mapper3.2.A.py|python reducer3.2.A.py|sort -k2,2n
```

```
reporter:counter:Reducer Counters,Calls,1
reporter:counter:Mapper Counters,Calls,1
bar      1
foo      1
labs     1
quux     1
quux     1
foo      2
```

```
In [13]: !chmod a+x mapper3.2.A.py
!chmod a+x reducer3.2.A.py
```

```
In [14]: # Hadoop command
# START STUDENT CODE HW32AHADOOP

!echo "foo foo quux labs foo bar quux" >foofoo.txt
!hdfs dfs -copyFromLocal foofoo.txt
!hdfs dfs -rm -r hw3.2.A-output

!hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyField
BasedComparator \
-D stream.num.map.output.key.fields=2 \
-D stream.map.output.field.separator="\t" \
-D mapreduce.partition.keycomparator.options="-k1,1nr -k2,2" \
-files mapper3.2.A.py,reducer3.2.A.py\
-mapper mapper3.2.A.py \
-reducer reducer3.2.A.py\
-input foofoo.txt \
-output hw3.2.A-output

# END STUDENT CODE HW32AHADOOP
```

```

copyFromLocal: `foofoo.txt': File exists
Deleted hw3.2.A-output
17/05/30 05:56:48 INFO Configuration.deprecation: mapred.output.key.comparator
.class is deprecated. Instead, use mapreduce.job.output.key.comparator.class
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob3826147749252105465.jar tmpDir=null
17/05/30 05:56:49 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:56:50 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:56:51 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 05:56:51 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 05:56:51 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0087
17/05/30 05:56:52 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0087
17/05/30 05:56:52 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0087/
17/05/30 05:56:52 INFO mapreduce.Job: Running job: job_1496033164706_0087
17/05/30 05:57:03 INFO mapreduce.Job: Job job_1496033164706_0087 running in ub
er mode : false
17/05/30 05:57:03 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 05:57:11 INFO mapreduce.Job: map 50% reduce 0%
17/05/30 05:57:13 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 05:57:20 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 05:57:20 INFO mapreduce.Job: Job job_1496033164706_0087 completed suc
cessfully
17/05/30 05:57:20 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=72
        FILE: Number of bytes written=353330
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=255
        HDFS: Number of bytes written=26
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=12831
        Total time spent by all reduces in occupied slots (ms)=4650
        Total time spent by all map tasks (ms)=12831
        Total time spent by all reduce tasks (ms)=4650
        Total vcore-seconds taken by all map tasks=12831
        Total vcore-seconds taken by all reduce tasks=4650
        Total megabyte-seconds taken by all map tasks=13138944
        Total megabyte-seconds taken by all reduce tasks=4761600
    Map-Reduce Framework
        Map input records=1
        Map output records=7
        Map output bytes=52
        Map output materialized bytes=78
        Input split bytes=208
        Combine input records=0
        Combine output records=0
        Reduce input groups=4
        Reduce shuffle bytes=78
        Reduce input records=7
        Reduce output records=4
        Spilled Records=14
        Shuffled Maps =2
        Failed Shuffles=0
        Merged Map outputs=2
        GC time elapsed (ms)=123
        CPU time spent (ms)=2640

```

In [15]:

```
#HDFS Output
!hdfs dfs -ls hw3.2.A-output
!hdfs dfs -cat hw3.2.A-output/part-0000*
```

Found 2 items
-rw-r--r-- 1 root supergroup 0 2017-05-30 05:57 hw3.2.A-output/_SUC
CESS
-rw-r--r-- 1 root supergroup 26 2017-05-30 05:57 hw3.2.A-output/part
-00000
quux 2
labs 1
foo 3
bar 1

INSERT SCREENSHOT OF JOB TRACKER UI COUNTERS

	Physical memory (bytes) snapshot	0	4	4
	Reduce input groups	0	7	7
	Reduce input records	0	4	4
	Reduce output records	0	78	78
	Reduce shuffle bytes	0	2	2
	Shuffled Maps	7	7	14
	Spilled Records	498597888	141033472	639631360
	Total committed heap usage (bytes)	2729283584	1351589888	4080873472
	Virtual memory (bytes) snapshot			
Mapper Counters	Name	Map	Reduce	Total
	Calls	2	0	2
Reducer Counters	Name	Map	Reduce	Total
	Calls	0	1	1
	Name	Map	Reduce	Total
	BAD_ID	0	0	0

3.2.A EXPLANATION

With default setting , MapReduce selected to partition records into two maps so mapper is called 2 times. And then results are sent to single reducer.

3.2.B SOLUTION


```
In [16]: %%writefile mapper3.2.B.py
#!/usr/bin/env python
# START STUDENT CODE HW32BMAPPER
from __future__ import division
import math
import os
import sys
import re

separator = ','
sys.stderr.write("reporter:counter:Mapper Counters,Calls,1\n")
WORD_RE = re.compile(r"[\w]+")
numReducers = int(os.environ.get('NUM_PARTITIONS', '4'))

def makeKey(word,n):
    divisor = 26/n
    return int(math.ceil((ord(word[0])-96)/divisor))

#loop through each records
for line in (sys.stdin):
    #get 3rd column
    fields = line.split(separator)
    if 'Complaint ID' != fields[0] :

        # we have a real record, so do some mapping
        counter_name = None
        for word in [s.lower() for s in WORD_RE.findall(fields[3])]:
            key = makeKey(word,numReducers)
            print '%s\t%s\t%s' % (key,word, 1)

# END STUDENT CODE HW32BMAPPER
```

Overwriting mapper3.2.B.py

```
In [17]: %%writefile reducer3.2.B.py
#!/usr/bin/env python
# START STUDENT CODE HW32BREDUCER
import sys

cur_key = None
cur_count = 0
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")
for line in sys.stdin:
    partkey,key, value = line.split()
    if key == cur_key:
        cur_count += int(value)
    else:
        if cur_key:
            print '%s\t%s' % (cur_key, cur_count)
        cur_key = key
        cur_count = int(value)

print '%s\t%s' % (cur_key, cur_count)
# END STUDENT CODE HW32BREDUCER
```

Overwriting reducer3.2.B.py

```
In [18]: !chmod a+x mapper3.2.B.py
!chmod a+x reducer3.2.B.py
```

```
In [19]: #unit test
!head -10 Consumer_Complaints.csv|python mapper3.2.B.py|sort -k1,1|python reducer3.2.B.py
```

```
reporter:counter:Reducer Counters,Calls,1
reporter:counter:Mapper Counters,Calls,1
and 1
attempts 1
club 1
collect 1
cont'd 1
credit 1
debt 3
deposits 1
disclosure 2
false 1
health 1
incorrect 1
information 1
lease 2
loan 2
managing 2
not 1
of 2
on 1
or 3
owed 1
report 1
representation 1
statements 1
the 2
verification 2
withdrawals 1
```

```
In [20]: # Hadoop command
# START STUDENT CODE HW32BHADOOP
!hdfs dfs -rm -r hw3.2.B-output
!hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
-D mapreduce.job.maps=2 \
-D mapreduce.job.reduces=2\
-D stream.num.map.output.key.fields=2 \
-D stream.map.output.field.separator="\t" \
-D mapreduce.partition.keypartitioner.options=-k1,1 \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyField
BasedComparator \
-D mapreduce.partition.keycomparator.options="-k1,1nr -k2,2" \
-files mapper3.2.B.py,reducer3.2.B.py\
-mapper mapper3.2.B.py \
-reducer reducer3.2.B.py\
-input Consumer_Complaints.csv \
-cmdenv NUM_PARTITIONS=2\
-output hw3.2.B-output

# END STUDENT CODE HW32BHADOOP
```

Deleted hw3.2.B-output

```
17/05/30 05:57:37 INFO Configuration.deprecation: mapred.output.key.comparator
.class is deprecated. Instead, use mapreduce.job.output.key.comparator.class
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob4182344244284091905.jar tmpDir=null
17/05/30 05:57:38 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:57:39 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:57:41 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 05:57:41 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 05:57:42 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0088
17/05/30 05:57:42 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0088
17/05/30 05:57:42 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0088/
17/05/30 05:57:42 INFO mapreduce.Job: Running job: job_1496033164706_0088
17/05/30 05:57:50 INFO mapreduce.Job: Job job_1496033164706_0088 running in ub
er mode : false
17/05/30 05:57:50 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 05:58:06 INFO mapreduce.Job: map 33% reduce 0%
17/05/30 05:58:07 INFO mapreduce.Job: map 67% reduce 0%
17/05/30 05:58:08 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 05:58:20 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 05:58:21 INFO mapreduce.Job: Job job_1496033164706_0088 completed suc
cessfully
17/05/30 05:58:21 INFO mapreduce.Job: Counters: 51
```

File System Counters

```
FILE: Number of bytes read=14174923
FILE: Number of bytes written=28821512
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=50910816
HDFS: Number of bytes written=2091
HDFS: Number of read operations=12
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
```

Job Counters

```
Launched map tasks=2
Launched reduce tasks=2
Data-local map tasks=2
Total time spent by all maps in occupied slots (ms)=29807
Total time spent by all reduces in occupied slots (ms)=17766
Total time spent by all map tasks (ms)=29807
Total time spent by all reduce tasks (ms)=17766
Total vcore-seconds taken by all map tasks=29807
Total vcore-seconds taken by all reduce tasks=17766
Total megabyte-seconds taken by all map tasks=30522368
Total megabyte-seconds taken by all reduce tasks=18192384
```

Map-Reduce Framework

```
Map input records=312913
Map output records=980482
Map output bytes=12213947
Map output materialized bytes=14174935
Input split bytes=234
Combine input records=0
Combine output records=0
Reduce input groups=169
Reduce shuffle bytes=14174935
Reduce input records=980482
Reduce output records=169
Spilled Records=1960964
Shuffled Maps =4
Failed Shuffles=0
Merged Map outputs=4
GC time elapsed (ms)=432
CPU time spent (ms)=15210
```

```
In [21]: # 3.2.B OUTPUT/ANSWER
!hdfs dfs -ls hw3.2.B-output
!echo "_____Output_____"
!hdfs dfs -cat hw3.2.B-output/part-0000*
```

Found 3 items		
-rw-r--r--	1 root supergroup	0 2017-05-30 05:58 hw3.2.B-output/_SUC
CESS		
-rw-r--r--	1 root supergroup	841 2017-05-30 05:58 hw3.2.B-output/part
-00000		
-rw-r--r--	1 root supergroup	1250 2017-05-30 05:58 hw3.2.B-output/part
-00001		
Output		
opening	16205	
other	7886	
out	1242	
pay	3821	
payment	92	
plans	350	
practices	1003	
privacy	240	
problems	9484	
rate	3431	
receiving	3226	
report	34903	
reporting	6559	
rewards	1002	
scam	566	
score	4357	
servicing	36767	
sharing	2832	
shopping	672	
statements	2508	
stop	131	
taking	3747	
transaction	1485	
underwriting	2774	
using	2422	
when	4095	
with	1944	
withdrawals	10555	
a	3503	
account	20681	
acct	163	
an	2505	
and	16448	
applied	139	
apr	3431	
arbitration	168	
available	274	
bankruptcy	222	
being	5663	
billing	8158	
by	5663	
can't	1999	
cash	240	
caused	5663	
changes	350	
charges	131	
checks	75	
closing	2795	
company's	4858	
cont'd	11848	
convenience	75	
credit	55251	
debt	19309	
delay	243	
delinquent	1061	
deposits	10555	
determination	1490	
did	139	
disclosure	5214	
disputes	6938	
expect	807	
false	2500	

INSERT SCREENSHOT OF JOB TRACKER UI COUNTERS

Virtual memory (bytes) snapshot		2720735232	2713747456	5434482688
Mapper Counters	Name ▲	Map ↕	Reduce ↕	Total ↕
	Calls	2	0	2
Reducer Counters	Name ▲	Map ↕	Reduce ↕	Total ↕
	Calls	0	2	2

3.2.C SOLUTION

```
In [22]: %%writefile mapper3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CMAPPER
from __future__ import division
import math
import os
import sys
import re

separator = ','
sys.stderr.write("reporter:counter:Mapper Counters,Calls,1\n")
WORD_RE = re.compile(r"[\w']+")
#numReducers = int(os.environ.get('NUM_PARTITIONS', '4'))

total = 0

def makeKey(word,n):
    divisor = 26/n
    return int(math.ceil((ord(word[0])-96)/divisor))

#loop through each records
for line in (sys.stdin):
    #get 3rd column
    fields = line.split(separator)
    if 'Complaint ID' != fields[0] :

        # we have a real record, so do some mapping
        counter_name = None
        for word in [s.lower() for s in WORD_RE.findall(fields[3])]:
            #key = makeKey(word,numReducers)
            print '%s\t%s' % (word, 1)
            total = total + 1
print '%s\t%s' % ("*total", total)

# END STUDENT CODE HW32CMAPPER
```

Overwriting mapper3.2.C.py

```
In [23]: %%writefile combiner3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CCOMBINER
import sys

cur_key = None
cur_count = 0
sys.stderr.write("reporter:counter:Combiner Counters,Calls,1\n")
for line in sys.stdin:
    key, value = line.split()
    if key == cur_key:
        cur_count += int(value)
    else:
        if cur_key:
            print '%s\t%s' % (cur_key, cur_count)
        cur_key = key
        cur_count = int(value)

print '%s\t%s' % (cur_key, cur_count)
# END STUDENT CODE HW32CCOMBINER
```

Overwriting combiner3.2.C.py

```
In [24]: %%writefile reducer3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CREducer
import sys

cur_key = None
cur_count = 0
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")
for line in sys.stdin:
    key, value = line.split()
    if key == cur_key:
        cur_count += int(value)
    else:
        if cur_key:
            print '%s\t%s' % (cur_key, cur_count)
        cur_key = key
        cur_count = int(value)

print '%s\t%s' % (cur_key, cur_count)
# END STUDENT CODE HW32CREducer
```

Overwriting reducer3.2.C.py

```
In [25]: !chmod a+x mapper3.2.C.py
!chmod a+x reducer3.2.C.py
!chmod a+x combiner3.2.C.py
```



```
In [26]: #unit Testing
!head -10 Consumer_Complaints.csv|python mapper3.2.C.py|sort -k1,1|python combi
ner3.2.C.py|python reducer3.2.C.py
```

```
reporter:counter:Reducer Counters,Calls,1
reporter:counter:Combiner Counters,Calls,1
reporter:counter:Mapper Counters,Calls,1
*total 38
and 1
attempts 1
club 1
collect 1
cont'd 1
credit 1
debt 3
deposits 1
disclosure 2
false 1
health 1
incorrect 1
information 1
lease 2
loan 2
managing 2
not 1
of 2
on 1
or 3
owed 1
report 1
representation 1
statements 1
the 2
verification 2
withdrawals 1
```

```
In [27]: # Hadoop command
# START STUDENT CODE HW32CHADOOP
!hdfs dfs -rm -r hw3.2.C-output
!hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
    -files mapper3.2.C.py,reducer3.2.C.py,combiner3.2.C.py\
    -mapper mapper3.2.C.py \
    -reducer reducer3.2.C.py\
    -combiner combiner3.2.C.py\
    -input Consumer_Complaints.csv \
    -output hw3.2.C-output \
    -numReduceTasks 4

# END STUDENT CODE HW32CHADOOP
```

```

Deleted hw3.2.C-output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob4293544407680665915.jar tmpDir=null
17/05/30 05:58:35 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:58:35 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:58:37 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 05:58:37 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 05:58:38 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0089
17/05/30 05:58:38 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0089
17/05/30 05:58:38 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0089/
17/05/30 05:58:38 INFO mapreduce.Job: Running job: job_1496033164706_0089
17/05/30 05:58:46 INFO mapreduce.Job: Job job_1496033164706_0089 running in ub
er mode : false
17/05/30 05:58:46 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 05:58:59 INFO mapreduce.Job: map 50% reduce 0%
17/05/30 05:59:00 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 05:59:10 INFO mapreduce.Job: map 100% reduce 25%
17/05/30 05:59:13 INFO mapreduce.Job: map 100% reduce 50%
17/05/30 05:59:14 INFO mapreduce.Job: map 100% reduce 75%
17/05/30 05:59:15 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 05:59:15 INFO mapreduce.Job: Job job_1496033164706_0089 completed suc
cessfully
17/05/30 05:59:15 INFO mapreduce.Job: Counters: 52
    File System Counters
        FILE: Number of bytes read=4488
        FILE: Number of bytes written=715820
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50910816
        HDFS: Number of bytes written=2105
        HDFS: Number of read operations=18
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=8
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=4
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=20694
        Total time spent by all reduces in occupied slots (ms)=40757
        Total time spent by all map tasks (ms)=20694
        Total time spent by all reduce tasks (ms)=40757
        Total vcore-seconds taken by all map tasks=20694
        Total vcore-seconds taken by all reduce tasks=40757
        Total megabyte-seconds taken by all map tasks=21190656
        Total megabyte-seconds taken by all reduce tasks=41735168
    Map-Reduce Framework
        Map input records=312913
        Map output records=980484
        Map output bytes=9272529
        Map output materialized bytes=4512
        Input split bytes=234
        Combine input records=980484
        Combine output records=315
        Reduce input groups=170
        Reduce shuffle bytes=4512
        Reduce input records=315
        Reduce output records=170
        Spilled Records=630
        Shuffled Maps =8
        Failed Shuffles=0
        Merged Map outputs=8
        GC time elapsed (ms)=444
        CPU time spent (ms)=11680

```

```
In [28]: # 3.2.C OUTPUT/ANSWER
!hdfs dfs -ls hw3.2.C-output
!echo "_____Output_____"
!hdfs dfs -cat hw3.2.C-output/part-0000*
```

Found 5 items			
-rw-r--r--	1	root supergroup	0 2017-05-30 05:59 hw3.2.C-output/_SUC
CESS			
-rw-r--r--	1	root supergroup	452 2017-05-30 05:59 hw3.2.C-output/part
-00000			
-rw-r--r--	1	root supergroup	600 2017-05-30 05:59 hw3.2.C-output/part
-00001			
-rw-r--r--	1	root supergroup	517 2017-05-30 05:59 hw3.2.C-output/part
-00002			
-rw-r--r--	1	root supergroup	536 2017-05-30 05:59 hw3.2.C-output/part
-00003			
Output			
a	3503		
account	20681		
acct	163		
applied	139		
available	274		
by	5663		
can't	1999		
cash	240		
caused	5663		
checks	75		
closing	2795		
company's	4858		
cont'd	11848		
debt	19309		
delinquent	1061		
disputes	6938		
for	929		
i	925		
incorrect	29133		
issuance	640		
issue	1098		
making	3226		
of	10885		
on	29069		
or	22533		
owed	11848		
payoff	1155		
processing	243		
repay	1647		
sale	139		
service	1518		
the	6248		
to	8401		
transfer	597		
unable	8178		
verification	5214		
was	274		
workout	350		
wrong	169		
your	3844		
advance	240		
amount	98		
apply	118		
atm	2422		
bank	202		
cancelling	2795		
card	4405		
collect	11848		
communication	6920		
costs	4350		
credited	92		
dealing	1944		
decision	2774		
didn't	925		
disclosures	64		
fee	3198		
funds	5663		
not	4257		

INSERT SCREENSHOT OF JOB TRACKER UI COUNTERS

Combiner Counters	Name	Map	Reduce	Total
	Calls	8	0	8
Mapper Counters	Name	Map	Reduce	Total
	Calls	2	0	2
Mapper Counters	Name	Map	Reduce	Total
	Calls	0	4	4

```
In [29]: %%writefile frequencies_mapper3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CFREQMAPPER
from __future__ import division
import math
import os
import sys
import re

separator = ','
sys.stderr.write("reporter:counter:Mapper Counters,Calls,1\n")
WORD_RE = re.compile(r"[\w']+")

#loop through each records
for line in sys.stdin:
    print line.strip()

# END STUDENT CODE HW32CFREQMAPPER
```

Overwriting frequencies_mapper3.2.C.py

```
In [30]: %%writefile frequencies_reducer3.2.C.py
#!/usr/bin/env python
# START STUDENT CODE HW32CFREQREDUCER
import sys

# Initialize variables
total = 0
cur_key = None
cur_count = 0
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")
for line in sys.stdin:

    fields = line.replace('\n','').split('\t')
    count = fields[1]
    word = fields[0]
    try:
        count = int(count)
    except ValueError:
        continue
    if word == '*total':
        total = total + float(count)
    else:
        print 's\t%s\t%.2f' % (word, count, float(count)/total)
        #print "{0:20}\t{1:10}\t{2}\n".format(word, count, float(count)/total)
# END STUDENT CODE HW32CFREQREDUCER
```

Overwriting frequencies_reducer3.2.C.py

```
In [31]: # Hadoop command
# START STUDENT CODE HW32CFREQHADOOP
!hdfs dfs -rm -r hw3.2.D-output
!hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
    -D stream.num.map.output.key.fields=4 \
    -D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyField
BasedComparator \
    -D mapreduce.partition.keycomparator.options="-k2,2nr -k1,1" \
    -D mapreduce.job.reduces=1 \
    -files frequencies_mapper3.2.C.py,frequencies_reducer3.2.C.py \
    -mapper frequencies_mapper3.2.C.py \
    -reducer frequencies_reducer3.2.C.py \
    -input hw3.2.C-output \
    -output hw3.2.D-output \
    -partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner
# END STUDENT CODE HW32CFREQHADOOP
```

```

Deleted hw3.2.D-output
17/05/30 05:59:27 INFO Configuration.deprecation: mapred.output.key.comparator
.class is deprecated. Instead, use mapreduce.job.output.key.comparator.class
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob8276092118841390770.jar tmpDir=null
17/05/30 05:59:28 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:59:28 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 05:59:29 INFO mapred.FileInputFormat: Total input paths to process :
4
17/05/30 05:59:29 INFO mapreduce.JobSubmitter: number of splits:4
17/05/30 05:59:29 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0090
17/05/30 05:59:30 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0090
17/05/30 05:59:30 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0090/
17/05/30 05:59:30 INFO mapreduce.Job: Running job: job_1496033164706_0090
17/05/30 05:59:38 INFO mapreduce.Job: Job job_1496033164706_0090 running in ub
er mode : false
17/05/30 05:59:38 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 05:59:47 INFO mapreduce.Job: map 25% reduce 0%
17/05/30 05:59:50 INFO mapreduce.Job: map 50% reduce 0%
17/05/30 05:59:52 INFO mapreduce.Job: map 75% reduce 0%
17/05/30 05:59:53 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 05:59:55 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 05:59:56 INFO mapreduce.Job: Job job_1496033164706_0090 completed suc
cessfully
17/05/30 05:59:56 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=2621
        FILE: Number of bytes written=594745
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=2581
        HDFS: Number of bytes written=3105
        HDFS: Number of read operations=15
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=4
        Launched reduce tasks=1
        Data-local map tasks=4
        Total time spent by all maps in occupied slots (ms)=37939
        Total time spent by all reduces in occupied slots (ms)=5807
        Total time spent by all map tasks (ms)=37939
        Total time spent by all reduce tasks (ms)=5807
        Total vcore-seconds taken by all map tasks=37939
        Total vcore-seconds taken by all reduce tasks=5807
        Total megabyte-seconds taken by all map tasks=38849536
        Total megabyte-seconds taken by all reduce tasks=5946368
    Map-Reduce Framework
        Map input records=170
        Map output records=170
        Map output bytes=2275
        Map output materialized bytes=2639
        Input split bytes=476
        Combine input records=0
        Combine output records=0
        Reduce input groups=170
        Reduce shuffle bytes=2639
        Reduce input records=170
        Reduce output records=169
        Spilled Records=340
        Shuffled Maps =4
        Failed Shuffles=0
        Merged Map outputs=4
        GC time elapsed (ms)=127

```



```
In [32]: # 3.2.C OUTPUT/ANSWER
!hdfs dfs -ls hw3.2.D-output
!echo "_____Output_____"

!hdfs dfs -cat hw3.2.D-output/part-0000*
```

Found 2 items				
-rw-r--r--	1	root	supergroup	0 2017-05-30 05:59 hw3.2.D-output/_SUC
CESS				
-rw-r--r--	1	root	supergroup	3105 2017-05-30 05:59 hw3.2.D-output/part
-00000				
Output				
loan	119630	0.122		
modification	70487	0.072		
credit	55251	0.056		
servicing	36767	0.037		
report	34903	0.036		
incorrect	29133	0.030		
information	29069	0.030		
on	29069	0.030		
or	22533	0.023		
account	20681	0.021		
debt	19309	0.020		
and	16448	0.017		
opening	16205	0.017		
club	12545	0.013		
health	12545	0.013		
not	12353	0.013		
attempts	11848	0.012		
collect	11848	0.012		
cont'd	11848	0.012		
owed	11848	0.012		
of	10885	0.011		
my	10731	0.011		
deposits	10555	0.011		
withdrawals	10555	0.011		
problems	9484	0.010		
application	8868	0.009		
to	8401	0.009		
unable	8178	0.008		
billing	8158	0.008		
other	7886	0.008		
disputes	6938	0.007		
communication	6920	0.007		
tactics	6920	0.007		
reporting	6559	0.007		
lease	6337	0.006		
the	6248	0.006		
being	5663	0.006		
by	5663	0.006		
caused	5663	0.006		
funds	5663	0.006		
low	5663	0.006		
process	5505	0.006		
disclosure	5214	0.005		
verification	5214	0.005		
managing	5006	0.005		
company's	4858	0.005		
investigation	4858	0.005		
identity	4729	0.005		
card	4405	0.004		
get	4357	0.004		
score	4357	0.004		
costs	4350	0.004		
settlement	4350	0.004		
improper	4309	0.004		
interest	4238	0.004		
protection	4139	0.004		
when	4095	0.004		
repaying	3844	0.004		
your	3844	0.004		
fraud	3842	0.004		
are	3821	0.004		
pay	3821	0.004		
you	3821	0.004		
taking	3747	0.004		

What are the top 50 most frequent terms in your word count analysis?

```
In [33]: !hdfs dfs -cat hw3.2.D-output/part-0000* | head -50
```

```

loan      119630  0.122
modification 70487  0.072
credit    55251  0.056
servicing 36767  0.037
report    34903  0.036
incorrect  29133  0.030
information 29069  0.030
on        29069  0.030
or        22533  0.023
account   20681  0.021
debt      19309  0.020
and       16448  0.017
opening   16205  0.017
club      12545  0.013
health    12545  0.013
not       12353  0.013
attempts  11848  0.012
collect   11848  0.012
cont'd    11848  0.012
owed      11848  0.012
of        10885  0.011
my        10731  0.011
deposits  10555  0.011
withdrawals 10555  0.011
problems  9484   0.010
application 8868  0.009
to        8401  0.009
unable    8178  0.008
billing   8158  0.008
other     7886  0.008
disputes  6938  0.007
communication 6920  0.007
tactics   6920  0.007
reporting 6559  0.007
lease     6337  0.006
the       6248  0.006
being     5663  0.006
by        5663  0.006
caused    5663  0.006
funds     5663  0.006
low       5663  0.006
process   5505  0.006
disclosure 5214  0.005
verification 5214  0.005
managing  5006  0.005
company's 4858  0.005
investigation 4858  0.005
identity  4729  0.005
card      4405  0.004
get       4357  0.004

```

Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order.

```
In [34]: !hdfs dfs -cat hw3.2.D-output/part-0000* | sort -k2,2nr | head -50
```

```
loan      119630  0.122
modification 70487  0.072
credit    55251  0.056
servicing 36767  0.037
report    34903  0.036
incorrect  29133  0.030
information 29069  0.030
on        29069  0.030
or        22533  0.023
account   20681  0.021
debt      19309  0.020
and       16448  0.017
opening   16205  0.017
club      12545  0.013
health    12545  0.013
not       12353  0.013
attempts  11848  0.012
collect   11848  0.012
cont'd    11848  0.012
owed      11848  0.012
of        10885  0.011
my        10731  0.011
deposits  10555  0.011
withdrawals 10555  0.011
problems  9484  0.010
application 8868  0.009
to        8401  0.009
unable    8178  0.008
billing   8158  0.008
other     7886  0.008
disputes  6938  0.007
communication 6920  0.007
tactics   6920  0.007
reporting 6559  0.007
lease     6337  0.006
the       6248  0.006
being     5663  0.006
by        5663  0.006
caused    5663  0.006
funds     5663  0.006
low       5663  0.006
process   5505  0.006
disclosure 5214  0.005
verification 5214  0.005
managing  5006  0.005
company's 4858  0.005
investigation 4858  0.005
identity  4729  0.005
card      4405  0.004
get       4357  0.004
```

Present bottom 10 tokens (least frequent items).

```
In [35]: !hdfs dfs -tail hw3.2.D-output/part-00000 > hw3.2.D.txt
!tail -10 hw3.2.D.txt
```

```
apply      118      0.000
amount     98      0.000
credited    92      0.000
payment    92      0.000
checks     75      0.000
convenience 75      0.000
amt        71      0.000
day        71      0.000
disclosures 64      0.000
missing    64      0.000
```

3.2.1

Using **2 reducers**: What are the top **50 most frequent terms** in your word count analysis?

Present the top 50 terms and their frequency and their relative frequency. Present the top 50 terms and their frequency and their relative frequency. If there are ties please sort the tokens in alphanumeric/string order. Present bottom 10 tokens (least frequent items). Please **use a combiner**.

START STUDENT CODE HW321 (INSERT CELLS BELOW AS NEEDED)

```
In [36]: %%writefile frequencies_mapper3.2.1.py
#!/usr/bin/env python
from __future__ import division
import math
import os
import sys
import re

count = 0

separator = ','
#create partition key
def makeKeyn(word):
    if ord(word[0]) in range(ord('a'), ord('m')):
        return 'A'
    else:
        return 'B'
#regex for word extraction
WORD_RE = re.compile(r"[\w']+")
for line in sys.stdin:
    fields = line.split(separator)
    if 'Complaint ID' != fields[0]:
        for word in [s.lower() for s in WORD_RE.findall(fields[3])]:
            # prepend a key based on the number of reducers
            key = makeKeyn(word)
            count = count + 1
            print key, "\t", word, "\t", 1
print 'A', "\t", "*total", "\t", count #to get total in all combiners
print 'B', "\t", "*total", "\t", count
```

Overwriting frequencies_mapper3.2.1.py

```
In [37]: %%writefile frequencies_reducer3.2.1.py
#!/usr/bin/env python
import sys
import os
# Initialize variables
total = 0
cur_key = None
cur_count = 0
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")
dictcounts = {}
#totalrecs = int(os.environ.get('TOTAL_RECS', '980482'))
for line in sys.stdin:

    fields = line.replace('\n','').split('\t')
    count = fields[2]
    word = fields[1]
    try:
        count = int(count)
    except ValueError:
        continue
    if word == '*total': #not required in multireducers
        total = total + int(count)
    else:
        x = dictcounts.get(word,None)
        if x != None:
            dictcounts[word]+=count
        else:
            dictcounts[word]=count
for key in dictcounts:

    print '%s\t%d\t%.3f' %(key,dictcounts[key] ,float(dictcounts[key])/total)

    #print '%s,%s\t%s\t%.3f' % (fields[0],word, count, float(count)/total)
```

Overwriting frequencies_reducer3.2.1.py

```
In [38]: %%writefile frequencies_combine3.2.1.py
#!/usr/bin/env python
import sys
import os
# Initialize variables
total = 0
cur_key = ("key1","key2")
cur_count = 0
sys.stderr.write("reporter:counter:combiner Counters,Calls,1\n")

for line in sys.stdin:

    partkey,key1, value = line.split()
    partkey = partkey
    key = (partkey,key1)
    if key1 == cur_key[1]:
        cur_count += int(value)
    else:
        if cur_key!= ("key1","key2"):
            print '%s\t%s\t%d' % (cur_key[0],cur_key[1], cur_count)
            cur_key = key
            cur_count = int(value)

print '%s\t%s\t%d' % (cur_key[0],cur_key[1], cur_count)
```

Overwriting frequencies_combine3.2.1.py

```
In [39]: !chmod a+x frequencies_reducer3.2.1.py
!chmod a+x frequencies_combine3.2.1.py
!chmod a+x frequencies_mapper3.2.1.py
```

END STUDENT CODE HW321

In [40]: *#start 3.2.1*

```
!hdfs dfs -rm -r hw3.2.1-output
!hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
  -D stream.num.map.output.key.fields=2 \
  -D mapreduce.partition.keypartitioner.options=-k1,1 \
  -D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.K
eyFieldBasedComparator \
  -D mapreduce.partition.keycomparator.options="-k1,1 -k2,2" \
  -D mapreduce.job.reduces=2 \
  -files frequencies_mapper3.2.1.py,frequencies_reducer3.2.1.py,frequencies_c
ombine3.2.1.py \
  -mapper frequencies_mapper3.2.1.py\
  -reducer frequencies_reducer3.2.1.py\
  -combiner frequencies_combine3.2.1.py\
  -input Consumer_Complaints.csv \
  -output hw3.2.1-output \
  -partitioner org.apache.hadoop.mapred.lib.KeyFieldBasedPartitioner
```



```

Deleted hw3.2.1-output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob5705506576480295317.jar tmpDir=null
17/05/30 06:00:21 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:00:21 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:00:22 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 06:00:22 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 06:00:23 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0091
17/05/30 06:00:23 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0091
17/05/30 06:00:23 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0091/
17/05/30 06:00:23 INFO mapreduce.Job: Running job: job_1496033164706_0091
17/05/30 06:00:31 INFO mapreduce.Job: Job job_1496033164706_0091 running in ub
er mode : false
17/05/30 06:00:31 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 06:00:46 INFO mapreduce.Job: map 28% reduce 0%
17/05/30 06:00:47 INFO mapreduce.Job: map 56% reduce 0%
17/05/30 06:00:49 INFO mapreduce.Job: map 62% reduce 0%
17/05/30 06:00:50 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 06:01:02 INFO mapreduce.Job: map 100% reduce 50%
17/05/30 06:01:03 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 06:01:04 INFO mapreduce.Job: Job job_1496033164706_0091 completed suc
cessfully
17/05/30 06:01:04 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=5142
        FILE: Number of bytes written=485598
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=50910816
        HDFS: Number of bytes written=3105
        HDFS: Number of read operations=12
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=4
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=2
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=32337
        Total time spent by all reduces in occupied slots (ms)=19546
        Total time spent by all map tasks (ms)=32337
        Total time spent by all reduce tasks (ms)=19546
        Total vcore-seconds taken by all map tasks=32337
        Total vcore-seconds taken by all reduce tasks=19546
        Total megabyte-seconds taken by all map tasks=33113088
        Total megabyte-seconds taken by all reduce tasks=20015104
    Map-Reduce Framework
        Map input records=312913
        Map output records=980486
        Map output bytes=13194501
        Map output materialized bytes=5154
        Input split bytes=234
        Combine input records=980486
        Combine output records=317
        Reduce input groups=2
        Reduce shuffle bytes=5154
        Reduce input records=317
        Reduce output records=169
        Spilled Records=634
        Shuffled Maps =4
        Failed Shuffles=0
        Merged Map outputs=4
        GC time elapsed (ms)=344
        CPU time spent (ms)=20020

```

```
In [41]: !hdfs dfs -ls hw3.2.1-output
!hdfs dfs -cat hw3.2.1-output/part-0000* | sort -k2,2nr > hw3.2.1.txt
```

```
Found 3 items
-rw-r--r--  1 root supergroup          0 2017-05-30 06:01 hw3.2.1-output/_SUCCESS
-rw-r--r--  1 root supergroup    1392 2017-05-30 06:01 hw3.2.1-output/part-00000
-rw-r--r--  1 root supergroup    1713 2017-05-30 06:01 hw3.2.1-output/part-00001
```

```
In [42]: !head -50 hw3.2.1.txt
```

```
loan      119630  0.122
modification  70487  0.072
credit    55251  0.056
servicing  36767  0.037
report    34903  0.036
incorrect  29133  0.030
information 29069  0.030
on        29069  0.030
or        22533  0.023
account   20681  0.021
debt      19309  0.020
and       16448  0.017
opening   16205  0.017
club      12545  0.013
health    12545  0.013
not       12353  0.013
attempts  11848  0.012
collect   11848  0.012
cont'd    11848  0.012
owed      11848  0.012
of        10885  0.011
my        10731  0.011
deposits  10555  0.011
withdrawals 10555  0.011
problems  9484   0.010
application 8868  0.009
to        8401  0.009
unable    8178  0.008
billing   8158  0.008
other     7886  0.008
disputes  6938  0.007
communication 6920  0.007
tactics   6920  0.007
reporting  6559  0.007
lease     6337  0.006
the       6248  0.006
being     5663  0.006
by        5663  0.006
caused    5663  0.006
funds     5663  0.006
low       5663  0.006
process   5505  0.006
disclosure 5214  0.005
verification 5214  0.005
managing  5006  0.005
company's  4858  0.005
investigation 4858  0.005
identity  4729  0.005
card      4405  0.004
get       4357  0.004
```

HW3.3. Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers.

For this homework use the online browsing behavior dataset located at:

<https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0>

Each line in this dataset represents a browsing session of a customer. On each line, each string of 8 characters represents the id of an item browsed during that session. The items are separated by spaces.

Here are the first few lines of the ProductPurchaseData FRO11987 ELE17451 ELE89019 SNA90258 GRO99222 GRO99222 GRO12298 FRO12685 ELE91550 SNA11465 ELE26917 ELE52966 FRO90334 SNA30755 ELE17451 FRO84225 SNA80192 ELE17451 GRO73461 DAI22896 SNA99873 FRO86643 ELE17451 ELE37798 FRO86643 GRO56989 ELE23393 SNA11465 ELE17451 SNA69641 FRO86643 FRO78087 SNA11465 GRO39357 ELE28573 ELE11375 DAI54444

Do some exploratory data analysis of this dataset guided by the following questions:

How many unique items are available from this supplier?

Using a single reducer: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

START STUDENT CODE HW33 (INSERT CELLS BELOW AS NEEDED)

```
In [43]: #get data first
!wget "https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0"
!mv ProductPurchaseData.txt?dl=0 ProductPurchaseData.txt

--2017-05-30 06:01:12-- https://www.dropbox.com/s/zlfyiwa70poqg74/ProductPurchaseData.txt?dl=0
Resolving www.dropbox.com... 162.125.4.1
Connecting to www.dropbox.com|162.125.4.1|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://dl.dropboxusercontent.com/content_link/YnjBvBDFNoHAIgv7R3QKC7QvPDOfghJjID5D4DL2UM0PKueav0KljF0dzVxrdSlZ/file [following]
--2017-05-30 06:01:13-- https://dl.dropboxusercontent.com/content_link/YnjBvBDFNoHAIgv7R3QKC7QvPDOfghJjID5D4DL2UM0PKueav0KljF0dzVxrdSlZ/file
Resolving dl.dropboxusercontent.com... 162.125.4.6
Connecting to dl.dropboxusercontent.com|162.125.4.6|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 3458517 (3.3M) [text/plain]
Saving to: `ProductPurchaseData.txt?dl=0'

100%[=====>] 3,458,517 2.46M/s in 1.3s

2017-05-30 06:01:15 (2.46 MB/s) - `ProductPurchaseData.txt?dl=0' saved [3458517/3458517]
```

```
In [44]: !hdfs dfs -copyFromLocal ProductPurchaseData.txt
!hdfs dfs -rm -r hw3.3-output

copyFromLocal: `ProductPurchaseData.txt': File exists
Deleted hw3.3-output
```

```
In [45]: %%writefile mapper_33.py
#!/usr/bin/python
## mapper.py

import sys

# Increment mapper counter
sys.stderr.write("reporter:counter:Mapper Counters,Calls,1\n")

# Initials
sys.stderr.write(ize variables
total = 0
basket_size = 0
largest_basket_size = 0

for line in sys.stdin:
    total = 0
    basket_size = 0
    # Split our line into products
    for product in line.replace('\n','').split():
        print '%s\t%s' % (product, 1)
        #print generateLongCountToken(product)
        basket_size += 1
        total += 1

    print '%s\t%s' % (*largest_basket', total)
    #basket_size = 0
```

Overwriting mapper_33.py

```
In [46]: !chmod a+x mapper_33.py
```

```

In [47]: %%writefile reducer33.py
#!/usr/bin/env python
# START STUDENT CODE HW32AREducer
import sys
from collections import OrderedDict

cur_key = None
cur_count = 0
dictcounts = {}
largest = []
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")
for line in sys.stdin:
    key, value = line.split()

    if key != '*largest_basket':
        if key == cur_key:
            cur_count += int(value)
        else:
            if cur_key:
                dictcounts[cur_key] = cur_count
                #print '%s\t%s' % (cur_key, cur_count)
                cur_key = key
                cur_count = int(value)
            else:
                if key == '*largest_basket':
                    largest.append(int(value))
print "*****Output*****"
print "Maximum length of Bucket %d"%(max(largest))

print "Total No of Unique products %d"% len(dictcounts.keys())
totals = sum(dictcounts.values())
dictcounts =OrderedDict(sorted(dictcounts.items(), key=lambda t: t[1], reverse=
True))
count = 0
print "*****Top 50 Products*****"
for key in dictcounts:
    if count <= 49:
        print '%s\t%d\t%.3f' %(key,dictcounts[key] ,float(dictcounts[key])/tot
als )
        count += 1

```

Overwriting reducer33.py

```
In [48]: !hdfs dfs -cat ProductPurchaseData.txt |python mapper_33.py|sort -k1,1 |python
reducer33.py
```

```
reporter:counter:Mapper Counters,Calls,1
reporter:counter:Reducer Counters,Calls,1
*****Output*****
Maximum length of Bucket 37
Total No of Unique products 12591
*****Top 50 Products*****
DAI62779      6667      0.018
FRO40251      3881      0.010
ELE17451      3875      0.010
GRO73461      3602      0.009
SNA80324      3044      0.008
ELE32164      2851      0.007
DAI75645      2736      0.007
SNA45677      2455      0.006
FRO31317      2330      0.006
DAI85309      2293      0.006
ELE26917      2292      0.006
FRO80039      2233      0.006
GRO21487      2115      0.006
SNA99873      2083      0.005
GRO59710      2004      0.005
GRO71621      1920      0.005
FRO85978      1918      0.005
GRO30386      1840      0.005
ELE74009      1816      0.005
GRO56726      1784      0.005
DAI63921      1773      0.005
GRO46854      1756      0.005
ELE66600      1713      0.004
DAI83733      1712      0.004
FRO32293      1702      0.004
ELE66810      1697      0.004
SNA55762      1646      0.004
DAI22177      1627      0.004
FRO78087      1531      0.004
ELE99737      1516      0.004
ELE34057      1489      0.004
GRO94758      1489      0.004
FRO35904      1436      0.004
FRO53271      1420      0.004
SNA93860      1407      0.004
SNA90094      1390      0.004
GRO38814      1352      0.004
ELE56788      1345      0.004
GRO61133      1321      0.003
ELE74482      1316      0.003
DAI88807      1316      0.003
ELE59935      1311      0.003
SNA96271      1295      0.003
DAI43223      1290      0.003
ELE91337      1289      0.003
GRO15017      1275      0.003
DAI31081      1261      0.003
GRO81087      1220      0.003
DAI22896      1219      0.003
GRO85051      1214      0.003
```

```
In [49]: !hdfs dfs -rm -r hw3.3-output

!hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
-D mapred.reduce.tasks=1 \
-D mapred.output.key.comparator.class=org.apache.hadoop.mapred.lib.KeyField
BasedComparator \
-D stream.num.map.output.key.fields=2 \
-D stream.map.output.field.separator="\t" \
-D mapreduce.partition.keycomparator.options="-k1,1" \
-files mapper_33.py, reducer33.py \
-mapper mapper_33.py \
-reducer reducer33.py \
-input ProductPurchaseData.txt \
-output hw3.3-output
```

```

rm: `hw3.3-output': No such file or directory
17/05/30 06:01:36 INFO Configuration.deprecation: mapred.reduce.tasks is depre
cated. Instead, use mapreduce.job.reduces
17/05/30 06:01:36 INFO Configuration.deprecation: mapred.output.key.comparator
.class is deprecated. Instead, use mapreduce.job.output.key.comparator.class
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob2739617999185734490.jar tmpDir=null
17/05/30 06:01:37 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:01:37 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:01:38 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 06:01:39 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 06:01:39 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0092
17/05/30 06:01:39 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0092
17/05/30 06:01:39 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0092/
17/05/30 06:01:39 INFO mapreduce.Job: Running job: job_1496033164706_0092
17/05/30 06:01:50 INFO mapreduce.Job: Job job_1496033164706_0092 running in ub
er mode : false
17/05/30 06:01:50 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 06:02:07 INFO mapreduce.Job: map 50% reduce 0%
17/05/30 06:02:08 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 06:02:17 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 06:02:18 INFO mapreduce.Job: Job job_1496033164706_0092 completed suc
cessfully
17/05/30 06:02:18 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=6005167
        FILE: Number of bytes written=12363340
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=3462847
        HDFS: Number of bytes written=1160
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=28944
        Total time spent by all reduces in occupied slots (ms)=7936
        Total time spent by all map tasks (ms)=28944
        Total time spent by all reduce tasks (ms)=7936
        Total vcore-seconds taken by all map tasks=28944
        Total vcore-seconds taken by all reduce tasks=7936
        Total megabyte-seconds taken by all map tasks=29638656
        Total megabyte-seconds taken by all reduce tasks=8126464
    Map-Reduce Framework
        Map input records=31101
        Map output records=411925
        Map output bytes=5181311
        Map output materialized bytes=6005173
        Input split bytes=234
        Combine input records=0
        Combine output records=0
        Reduce input groups=41578
        Reduce shuffle bytes=6005173
        Reduce input records=411925
        Reduce output records=54
        Spilled Records=823850
        Shuffled Maps =2
        Failed Shuffles=0
        Merged Map outputs=2
        GC time elapsed (ms)=126

```



```
In [50]: !hdfs dfs -ls hw3.3-output
!hdfs dfs -cat hw3.3-output/part-0000*
```

```
Found 2 items
-rw-r--r-- 1 root supergroup          0 2017-05-30 06:02 hw3.3-output/_SUCCESS
-rw-r--r-- 1 root supergroup    1160 2017-05-30 06:02 hw3.3-output/part-0000
*****Output*****
Maximum length of Bucket 37
Total No of Unique products 12591
*****Top 50 Products*****
DAI62779      6667      0.018
FRO40251      3881      0.010
ELE17451      3875      0.010
GRO73461      3602      0.009
SNA80324      3044      0.008
ELE32164      2851      0.007
DAI75645      2736      0.007
SNA45677      2455      0.006
FRO31317      2330      0.006
DAI85309      2293      0.006
ELE26917      2292      0.006
FRO80039      2233      0.006
GRO21487      2115      0.006
SNA99873      2083      0.005
GRO59710      2004      0.005
GRO71621      1920      0.005
FRO85978      1918      0.005
GRO30386      1840      0.005
ELE74009      1816      0.005
GRO56726      1784      0.005
DAI63921      1773      0.005
GRO46854      1756      0.005
ELE66600      1713      0.004
DAI83733      1712      0.004
FRO32293      1702      0.004
ELE66810      1697      0.004
SNA55762      1646      0.004
DAI22177      1627      0.004
FRO78087      1531      0.004
ELE99737      1516      0.004
ELE34057      1489      0.004
GRO94758      1489      0.004
FRO35904      1436      0.004
FRO53271      1420      0.004
SNA93860      1407      0.004
SNA90094      1390      0.004
GRO38814      1352      0.004
ELE56788      1345      0.004
GRO61133      1321      0.003
ELE74482      1316      0.003
DAI88807      1316      0.003
ELE59935      1311      0.003
SNA96271      1295      0.003
DAI43223      1290      0.003
ELE91337      1289      0.003
GRO15017      1275      0.003
DAI31081      1261      0.003
GRO81087      1220      0.003
DAI22896      1219      0.003
GRO85051      1214      0.003
```

END STUDENT CODE HW33

HW3.3.1 OPTIONAL

Using 2 reducers: Report your findings such as number of unique products; largest basket; report the top 50 most frequently purchased items, their frequency, and their relative frequency (break ties by sorting the products alphabetical order) etc. using Hadoop Map-Reduce.

START STUDENT CODE HW331 (INSERT CELLS BELOW AS NEEDED)

END STUDENT CODE HW331

HW3.4. (Computationally prohibitive but then again Hadoop can handle this) Pairs

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a map-reduce program to find products which are frequently browsed together. Fix the support count (cooccurrence count) to $s = 100$ (i.e. product pairs need to occur together at least 100 times to be considered frequent) and find pairs of items (sometimes referred to itemsets of size 2 in association rule mining) that have a support count of 100 or more.

List the top 50 product pairs with corresponding support count (aka frequency), and relative frequency or support (number of records where they occur, the number of records where they occur/the number of baskets in the dataset) in decreasing order of support for frequent ($100 > \text{count}$) itemsets of size 2.

Use the Pairs pattern (lecture 3) to extract these frequent itemsets of size 2. Free free to use combiners if they bring value. Instrument your code with counters for count the number of times your mapper, combiner and reducers are called.

Please output records of the following form for the top 50 pairs (itemsets of size 2):

```
item1, item2, support count, support
```

Fix the ordering of the pairs lexicographically (left to right), and break ties in support (between pairs, if any exist) by taking the first ones in lexicographically increasing order.

Report the compute time for the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers) Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts.

START STUDENT CODE HW34 (INSERT CELLS BELOW AS NEEDED)

```
In [51]: %%writefile mapper_34.py
#!/usr/bin/python
## mapper.py

import sys
from itertools import combinations

# Increment mapper counter
sys.stderr.write("reporter:counter:Mapper Counters,Calls,1\n")

# Initialize variables
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Split our line into products
    products = line.replace('\n','').split()

    # Get all combinations of products:
    # - Use a set to remove duplicate products
    # - Combinations finds tuples of length 2 with no repeats
    for pair in combinations(sorted(set(products)), 2):
        print '%s\t%s\t%s' % (pair[0], pair[1], 1)

    total += 1

# Print total words
print '%s\t%s\t%s' % ('*total', '*total', total)
```

Overwriting mapper_34.py

```
In [52]: !chmod a+x mapper_34.py
```

```
In [53]: #unit test
!hdfs dfs -cat ProductPurchaseData.txt |head -1|python mapper_34.py|sort -k1,1
```

```
reporter:counter:Mapper Counters,Calls,1
cat: Unable to write to output stream.
*total *total 1
ELE17451      ELE89019      1
ELE17451      FRO11987      1
ELE17451      GRO99222      1
ELE17451      SNA90258      1
ELE89019      FRO11987      1
ELE89019      GRO99222      1
ELE89019      SNA90258      1
FRO11987      GRO99222      1
FRO11987      SNA90258      1
GRO99222      SNA90258      1
```

```
In [54]: %%writefile combiner34.py
#!/usr/bin/env python
# START STUDENT CODE HW32CCOMBINER
import sys

cur_key = None
cur_count = 0
mydict = {}
sys.stderr.write("reporter:counter:Combiner Counters,Calls,1\n")
for line in sys.stdin:
    key1,key2, value = line.split()
    key = (key1,key2)
    if key == cur_key:
        cur_count += int(value)
    else:
        if cur_key and cur_count >=100:
            print '%s\t%s\t%s' % (cur_key[0],cur_key[1], cur_count)
            cur_key = key
            cur_count = int(value)

print '%s\t%s\t%s' % (cur_key[0],cur_key[1], cur_count)
```

Overwriting combiner34.py

```
In [55]: !chmod a+x combiner34.py
```

```
In [56]: %%writefile reducer34.py
#!/usr/bin/env python
# START STUDENT CODE HW32AREducer
import sys
from collections import OrderedDict

cur_key = None
cur_count = 0
dictcounts = {}
largest = []
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")
for line in sys.stdin:
    key1,key2, value = line.split()
    key = (key1,key2)
    if key1 != '*total':
        if key == cur_key:
            cur_count += int(value)
        else:
            if cur_key:
                dictcounts[cur_key] = cur_count
                #print '%s\t%s' % (cur_key, cur_count)
                cur_key = key
                cur_count = int(value)
            else:
                if key1 == '*total':
                    largest.append(int(value))

totals = sum(largest)
dictcounts =OrderedDict(sorted(dictcounts.items(), key=lambda t: t[1], reverse=
True))
count = 0
print "*****Top 50 Products*****"
for key in dictcounts:
    if count <= 49:
        print '%s\t%s\t%d\t%.3f' % (key[0],key[1],dictcounts[key] ,float(dictco
unts[key])/totals )
        count += 1
```

Overwriting reducer34.py

```
In [57]: !chmod a+x reducer34.py
```

```
In [58]: #hadoop call
!hdfs dfs -rm -r hw3.4-output
!time hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
-D stream.num.map.output.key.fields=4 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.K
eyFieldBasedComparator \
-D mapreduce.partition.keycomparator.options="-k1,1 -k2,2" \
-files mapper_34.py,combiner34.py,reducer34.py\
-mapper mapper_34.py\
-reducer reducer34.py\
-input ProductPurchaseData.txt \
-output hw3.4-output \
```

```

Deleted hw3.4-output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob7922580423188743914.jar tmpDir=null
17/05/30 06:02:42 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:02:42 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:02:44 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 06:02:44 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 06:02:44 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0093
17/05/30 06:02:45 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0093
17/05/30 06:02:45 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0093/
17/05/30 06:02:45 INFO mapreduce.Job: Running job: job_1496033164706_0093
17/05/30 06:03:02 INFO mapreduce.Job: Job job_1496033164706_0093 running in ub
er mode : false
17/05/30 06:03:02 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 06:03:20 INFO mapreduce.Job: map 20% reduce 0%
17/05/30 06:03:22 INFO mapreduce.Job: map 38% reduce 0%
17/05/30 06:03:23 INFO mapreduce.Job: map 51% reduce 0%
17/05/30 06:03:25 INFO mapreduce.Job: map 59% reduce 0%
17/05/30 06:03:28 INFO mapreduce.Job: map 67% reduce 0%
17/05/30 06:03:34 INFO mapreduce.Job: map 83% reduce 0%
17/05/30 06:03:36 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 06:03:48 INFO mapreduce.Job: map 100% reduce 76%
17/05/30 06:03:51 INFO mapreduce.Job: map 100% reduce 88%
17/05/30 06:03:54 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 06:04:12 INFO mapreduce.Job: Job job_1496033164706_0093 completed suc
cessfully
17/05/30 06:04:12 INFO mapreduce.Job: Counters: 51
  File System Counters
    FILE: Number of bytes read=58282374
    FILE: Number of bytes written=116918585
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=3462847
    HDFS: Number of bytes written=1442
    HDFS: Number of read operations=9
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=2
    Launched reduce tasks=1
    Data-local map tasks=2
    Total time spent by all maps in occupied slots (ms)=61869
    Total time spent by all reduces in occupied slots (ms)=34671
    Total time spent by all map tasks (ms)=61869
    Total time spent by all reduce tasks (ms)=34671
    Total vcore-seconds taken by all map tasks=61869
    Total vcore-seconds taken by all reduce tasks=34671
    Total megabyte-seconds taken by all map tasks=63353856
    Total megabyte-seconds taken by all reduce tasks=35503104
  Map-Reduce Framework
    Map input records=31101
    Map output records=2534016
    Map output bytes=53214336
    Map output materialized bytes=58282380
    Input split bytes=234
    Combine input records=0
    Combine output records=0
    Reduce input groups=877097
    Reduce shuffle bytes=58282380
    Reduce input records=2534016
    Reduce output records=51
    Spilled Records=5068032
    Shuffled Map = 2

```

```
In [59]: #print output
!hdfs dfs -ls hw3.4-output
!hdfs dfs -cat hw3.4-output/part-0000*
```

```
Found 2 items
-rw-r--r--  1 root supergroup          0 2017-05-30 06:04 hw3.4-output/_SUCCESS
SS
-rw-r--r--  1 root supergroup      1442 2017-05-30 06:04 hw3.4-output/part-0000
*****Top 50 Products*****
DAI62779      ELE17451      1592      0.051
FRO40251      SNA80324      1412      0.045
DAI75645      FRO40251      1254      0.040
FRO40251      GRO85051      1213      0.039
DAI62779      GRO73461      1139      0.037
DAI75645      SNA80324      1130      0.036
DAI62779      FRO40251      1070      0.034
DAI62779      SNA80324      923       0.030
DAI62779      DAI85309      918       0.030
ELE32164      GRO59710      911       0.029
FRO40251      GRO73461      882       0.028
DAI62779      DAI75645      882       0.028
DAI62779      ELE92920      877       0.028
FRO40251      FRO92469      835       0.027
DAI62779      ELE32164      832       0.027
DAI75645      GRO73461      712       0.023
DAI43223      ELE32164      711       0.023
DAI62779      GRO30386      709       0.023
ELE17451      FRO40251      697       0.022
DAI85309      ELE99737      659       0.021
DAI62779      ELE26917      650       0.021
GRO21487      GRO73461      631       0.020
DAI62779      SNA45677      604       0.019
ELE17451      SNA80324      597       0.019
DAI62779      GRO71621      595       0.019
DAI62779      SNA55762      593       0.019
DAI62779      DAI83733      586       0.019
ELE17451      GRO73461      580       0.019
GRO73461      SNA80324      562       0.018
DAI62779      GRO59710      561       0.018
DAI62779      FRO80039      550       0.018
DAI75645      ELE17451      547       0.018
DAI62779      SNA93860      537       0.017
DAI55148      DAI62779      526       0.017
DAI43223      GRO59710      512       0.016
ELE17451      ELE32164      511       0.016
DAI62779      SNA18336      506       0.016
ELE32164      GRO73461      486       0.016
DAI85309      ELE17451      482       0.015
DAI62779      FRO78087      482       0.015
DAI62779      GRO94758      479       0.015
DAI62779      GRO21487      471       0.015
GRO85051      SNA80324      471       0.015
ELE17451      GRO30386      468       0.015
FRO85978      SNA95666      463       0.015
DAI62779      FRO19221      462       0.015
DAI62779      GRO46854      461       0.015
DAI43223      DAI62779      459       0.015
ELE92920      SNA18336      455       0.015
DAI88079      FRO40251      446       0.014
```

END STUDENT CODE HW34

HW3.5: Stripes

Repeat 3.4 using the stripes design pattern for finding cooccurring pairs.

Report the compute times for stripes job versus the Pairs job. Describe the computational setup used (E.g., single computer; dual core; linux, number of mappers, number of reducers)

Instrument your mapper, combiner, and reducer to count how many times each is called using Counters and report these counts. Discuss the differences in these counts between the Pairs and Stripes jobs

START STUDENT CODE HW35 (INSERT CELLS BELOW AS NEEDED)

```
In [60]: %%writefile mapper_35.py
#!/usr/bin/python
## mapper.py

import sys
from itertools import combinations
#import collections

# Increment mapper counter
sys.stderr.write("reporter:counter:Mapper Counters,Calls,1\n")

# Initialize variables
total = 0

# Our input comes from STDIN (standard input)
for line in sys.stdin:
    # Split our line into products
    products = line.replace('\n','').split()

    # Get all combinations of products:
    # - Use a set to remove duplicate products
    # - Combinations finds tuples of length 2 with no repeats

    for i, term in enumerate(products):
        # Create a new stripe for each term
        stripe = {}

        for j, token in enumerate(products):
            # Don't count the term's co-occurrence with itself
            if i != j:
                x = stripe.get(token, None)
                if x == None:
                    stripe[token] = 1
                else:
                    stripe[token] += 1

        # Emit the term and the stripe
        print '%s\t%s' % (term, stripe)
# Increment total number of baskets
total += 1
stripe = {}
stripe['*total'] = total
print '%s\t%s' % (*total, stripe)
```

Overwriting mapper_35.py


```

In [61]: %%writefile reducer35.py
#!/usr/bin/env python
# START STUDENT CODE HW32AREducer
import sys
from collections import OrderedDict
#from collections import collections
#import collections
prev_key = None
cur_count = 0
prev_stripe = {}
largest = []
dictcounts = {}
sys.stderr.write("reporter:counter:Reducer Counters,Calls,1\n")

for line in sys.stdin:

    fields = line.replace('\n','').split('\t')
    key = fields[0]

    stripe = eval(fields[1])

    if prev_key == key:
        # We need to move through the dictionary and update counts
        for item in stripe:
            if item in prev_stripe:
                prev_stripe[item] += stripe[item]
            else:
                prev_stripe[item] = stripe[item]

    else:
        if len(prev_stripe) > 0:
            # We are at a new pair, need to print previous pair sum
            #print '%s\t%s' % (prev_key, prev_stripe)
            for word in prev_stripe:
                dictcounts[(prev_key,word)] = prev_stripe[word]
            prev_stripe = stripe
            prev_key = key

# Output the last line
if prev_stripe == stripe:
    for word in prev_stripe:
        dictcounts[(prev_key,word)] = prev_stripe[word]
totals = dictcounts[('*total','*total')]
dictcounts =OrderedDict(sorted(dictcounts.items(), key=lambda t: t[1], reverse=
True))
count = 0
print "*****Top 50 Products*****"
for key in dictcounts:
    if count <= 100 and count%2 == 0:
        print '%s\t%s\t%d\t%.3f' %(key[0],key[1],dictcounts[key] ,float(dictco
unts[key])/totals )
        count += 1

Overwriting reducer35.py

```

```

In [62]: !chmod a+x mapper_35.py
!chmod a+x reducer35.py

```

```
In [63]: !hdfs dfs -cat ProductPurchaseData.txt |head -10|python mapper_35.py|sort -k1,1
|python reducer35.py

reporter:counter:Mapper Counters,Calls,1
reporter:counter:Reducer Counters,Calls,1
cat: Unable to write to output stream.
*****Top 50 Products*****
*total *total 10 1.000
ELE17451 SNA80192 5 0.500
SNA69641 ELE17451 3 0.300
ELE17451 SNA69641 3 0.300
GRO73461 DAI22896 3 0.300
GRO73461 ELE17451 3 0.300
ELE17451 SNA11465 3 0.300
ELE17451 GRO73461 3 0.300
SNA90258 ELE17451 3 0.300
DAI22896 GRO73461 3 0.300
ELE17451 FRO86643 3 0.300
SNA11465 FRO86643 2 0.200
FRO78087 ELE11375 2 0.200
SNA69641 FRO78087 2 0.200
ELE17451 SNA85662 2 0.200
SNA80192 FRO18919 2 0.200
GRO73461 SNA99873 2 0.200
FRO81176 GRO94758 2 0.200
DAI91535 GRO94758 2 0.200
GRO56989 ELE37798 2 0.200
ELE17451 ELE37798 2 0.200
GRO99222 SNA80192 2 0.200
ELE28573 SNA69641 2 0.200
DAI22896 SNA80192 2 0.200
SNA69641 ELE28573 2 0.200
DAI22177 ELE17451 2 0.200
DAI91535 FRO81176 2 0.200
GRO94758 ELE17451 2 0.200
ELE23393 ELE17451 2 0.200
SNA85662 ELE17451 2 0.200
ELE17451 FRO81176 2 0.200
SNA80192 GRO94758 2 0.200
DAI91535 SNA80192 2 0.200
FRO86643 SNA11465 2 0.200
DAI22177 SNA85662 2 0.200
ELE28573 ELE11375 2 0.200
SNA80192 DAI91535 2 0.200
ELE28573 ELE17451 2 0.200
GRO94758 SNA80192 2 0.200
ELE11375 SNA69641 2 0.200
SNA85662 SNA80192 2 0.200
GRO94758 FRO81176 2 0.200
SNA80192 DAI22896 2 0.200
DAI22896 SNA99873 2 0.200
ELE66810 SNA80192 1 0.100
GRO73461 DAI91535 1 0.100
ELE59935 FRO18919 1 0.100
GRO39357 DAI54444 1 0.100
FRO84225 FRO90334 1 0.100
ELE17451 ELE89019 1 0.100
GRO75578 ELE17451 1 0.100
```

```
In [64]: #hadoop call
!hdfs dfs -rm -r hw3.5-output
!time hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar \
-D stream.num.map.output.key.fields=4 \
-D mapreduce.job.reduces=1 \
-D mapreduce.job.output.key.comparator.class=org.apache.hadoop.mapred.lib.K
eyFieldBasedComparator \
-D mapreduce.partition.keycomparator.options="-k1,1 -k2,2" \
-files mapper_35.py,combiner34.py,Reducer35.py\
-mapper mapper_35.py\
-reducer reducer35.py\
-input ProductPurchaseData.txt \
-output hw3.5-output \
-cmdenv PATH=/opt/anaconda/bin:$PATH
```

```

Deleted hw3.5-output
packageJobJar: [] [/usr/jars/hadoop-streaming-2.6.0-cdh5.7.0.jar] /tmp/streamj
ob6862122422670722550.jar tmpDir=null
17/05/30 06:04:40 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:04:40 INFO client.RMPProxy: Connecting to ResourceManager at /0.0.0
.0:8032
17/05/30 06:04:43 INFO mapred.FileInputFormat: Total input paths to process :
1
17/05/30 06:04:43 INFO mapreduce.JobSubmitter: number of splits:2
17/05/30 06:04:44 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_
1496033164706_0094
17/05/30 06:04:44 INFO impl.YarnClientImpl: Submitted application application_
1496033164706_0094
17/05/30 06:04:44 INFO mapreduce.Job: The url to track the job: http://quickst
art.cloudera:8088/proxy/application_1496033164706_0094/
17/05/30 06:04:44 INFO mapreduce.Job: Running job: job_1496033164706_0094
17/05/30 06:04:54 INFO mapreduce.Job: Job job_1496033164706_0094 running in ub
er mode : false
17/05/30 06:04:54 INFO mapreduce.Job: map 0% reduce 0%
17/05/30 06:05:13 INFO mapreduce.Job: map 58% reduce 0%
17/05/30 06:05:16 INFO mapreduce.Job: map 67% reduce 0%
17/05/30 06:05:22 INFO mapreduce.Job: map 83% reduce 0%
17/05/30 06:05:23 INFO mapreduce.Job: map 100% reduce 0%
17/05/30 06:05:36 INFO mapreduce.Job: map 100% reduce 71%
17/05/30 06:05:40 INFO mapreduce.Job: map 100% reduce 75%
17/05/30 06:05:43 INFO mapreduce.Job: map 100% reduce 78%
17/05/30 06:05:46 INFO mapreduce.Job: map 100% reduce 82%
17/05/30 06:05:49 INFO mapreduce.Job: map 100% reduce 86%
17/05/30 06:05:52 INFO mapreduce.Job: map 100% reduce 89%
17/05/30 06:05:55 INFO mapreduce.Job: map 100% reduce 92%
17/05/30 06:05:58 INFO mapreduce.Job: map 100% reduce 96%
17/05/30 06:06:01 INFO mapreduce.Job: map 100% reduce 99%
17/05/30 06:06:04 INFO mapreduce.Job: map 100% reduce 100%
17/05/30 06:06:19 INFO mapreduce.Job: Job job_1496033164706_0094 completed suc
cessfully
17/05/30 06:06:19 INFO mapreduce.Job: Counters: 51
    File System Counters
        FILE: Number of bytes read=81828248
        FILE: Number of bytes written=164010753
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=3462847
        HDFS: Number of bytes written=1468
        HDFS: Number of read operations=9
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=2
        Launched reduce tasks=1
        Data-local map tasks=2
        Total time spent by all maps in occupied slots (ms)=51492
        Total time spent by all reduces in occupied slots (ms)=52635
        Total time spent by all map tasks (ms)=51492
        Total time spent by all reduce tasks (ms)=52635
        Total vcore-seconds taken by all map tasks=51492
        Total vcore-seconds taken by all reduce tasks=52635
        Total megabyte-seconds taken by all map tasks=52727808
        Total megabyte-seconds taken by all reduce tasks=53898240
    Map-Reduce Framework
        Map input records=31101
        Map output records=380826
        Map output bytes=80638408
        Map output materialized bytes=81828254
        Input split bytes=234
        Combine input records=0
        Combine output records=0
        Reduce input groups=377574
        Reduce shuffle bytes=81828254

```

```
In [65]: #print output
!hdfs dfs -ls hw3.5-output
!hdfs dfs -cat hw3.5-output/part-0000*
```

```
Found 2 items
-rw-r--r-- 1 root supergroup          0 2017-05-30 06:06 hw3.5-output/_SUCCESS
SS
-rw-r--r-- 1 root supergroup    1468 2017-05-30 06:06 hw3.5-output/part-0000
*****Top 50 Products*****
*total *total 31101 1.000
ELE17451      DAI62779      1592    0.051
FRO40251      SNA80324      1412    0.045
FRO40251      DAI75645      1254    0.040
FRO40251      GRO85051      1213    0.039
GRO73461      DAI62779      1139    0.037
DAI75645      SNA80324      1130    0.036
FRO40251      DAI62779      1070    0.034
SNA80324      DAI62779      923     0.030
DAI85309      DAI62779      918     0.030
GRO59710      ELE32164      911     0.029
FRO40251      GRO73461      882     0.028
DAI62779      DAI75645      882     0.028
ELE92920      DAI62779      877     0.028
FRO92469      FRO40251      835     0.027
DAI62779      ELE32164      832     0.027
DAI75645      GRO73461      712     0.023
DAI43223      ELE32164      711     0.023
GRO30386      DAI62779      709     0.023
FRO40251      ELE17451      697     0.022
ELE99737      DAI85309      659     0.021
ELE26917      DAI62779      650     0.021
GRO73461      GRO21487      631     0.020
DAI62779      SNA45677      604     0.019
SNA80324      ELE17451      597     0.019
GRO71621      DAI62779      595     0.019
DAI62779      SNA55762      593     0.019
DAI62779      DAI83733      586     0.019
GRO73461      ELE17451      580     0.019
SNA80324      GRO73461      562     0.018
GRO59710      DAI62779      561     0.018
FRO80039      DAI62779      550     0.018
ELE17451      DAI75645      547     0.018
DAI62779      SNA93860      537     0.017
DAI55148      DAI62779      526     0.017
GRO59710      DAI43223      512     0.016
ELE17451      ELE32164      511     0.016
SNA18336      DAI62779      506     0.016
GRO73461      ELE32164      486     0.016
DAI62779      FRO78087      482     0.015
FRO78087      DAI62779      482     0.015
GRO94758      DAI62779      479     0.015
DAI62779      GRO21487      471     0.015
GRO21487      DAI62779      471     0.015
ELE17451      GRO30386      468     0.015
FRO85978      SNA95666      463     0.015
DAI62779      FRO19221      462     0.015
DAI62779      GRO46854      461     0.015
DAI62779      DAI43223      459     0.015
SNA18336      ELE92920      455     0.015
FRO40251      DAI88079      446     0.014
```

```
In [66]: !cat /proc/cpuinfo | grep processor | wc -l
```

2

```
In [67]: !cat /proc/meminfo
```

```
MemTotal:      5068464 kB
MemFree:       1719108 kB
MemAvailable:  1947436 kB
Buffers:       77720  kB
Cached:        477608 kB
SwapCached:    808    kB
Active:        2839068 kB
Inactive:      369844 kB
Active(anon):  2597472 kB
Inactive(anon): 204776 kB
Active(file):  241596 kB
Inactive(file): 165068 kB
Unevictable:   0      kB
Mlocked:       0      kB
SwapTotal:     1048572 kB
SwapFree:      1040120 kB
Dirty:         764    kB
Writeback:     0      kB
AnonPages:     2652740 kB
Mapped:        124824 kB
Shmem:         148664 kB
Slab:          87636  kB
SReclaimable:  61824  kB
SUnreclaim:    25812  kB
KernelStack:   15392  kB
PageTables:    12984  kB
NFS_Unstable:  0      kB
Bounce:        0      kB
WritebackTmp:  0      kB
CommitLimit:   3582804 kB
Committed_AS:  6061972 kB
VmallocTotal:  34359738367 kB
VmallocUsed:    0      kB
VmallocChunk:   0      kB
AnonHugePages:  0      kB
ShmemHugePages: 0      kB
ShmemPmdMapped: 0      kB
HugePages_Total: 0
HugePages_Free:  0
HugePages_Rsvd:  0
HugePages_Surp:  0
Hugepagesize:   2048  kB
DirectMap4k:    24576  kB
DirectMap2M:    4169728 kB
DirectMap1G:    3145728 kB
```


Answer

System Setup

Single Computer , docker Container, 2 Cores and 5GB RAM.

How many times is each mapper and reducer called?

Mapper 2 Reducer 1

Total time

With Pairs

```
real    1m6.364s
user    0m4.720s
sys     0m1.240s
```

```
Launched map tasks=2
Launched reduce tasks=1
Data-local map tasks=2
```

Total time spent by all maps in occupied slots (ms)=39445

```
Total time spent by all reduces in occupied slots (ms)=27979
Total time spent by all map tasks (ms)=39445
Total time spent by all reduce tasks (ms)=27979
Total vcore-seconds taken by all map tasks=39445
Total vcore-seconds taken by all reduce tasks=27979
Total megabyte-seconds taken by all map tasks=40391680
Total megabyte-seconds taken by all reduce tasks=28650496
```

With Stripes

```
real    1m11.730s
user    0m5.310s
sys     0m1.310s
```

```
Launched map tasks=2
Launched reduce tasks=1
Data-local map tasks=2
```

Total time spent by all maps in occupied slots (ms)=26607

```
Total time spent by all reduces in occupied slots (ms)=39085
Total time spent by all map tasks (ms)=26607
Total time spent by all reduce tasks (ms)=39085
Total vcore-seconds taken by all map tasks=26607
Total vcore-seconds taken by all reduce tasks=39085
Total megabyte-seconds taken by all map tasks=27245568
Total megabyte-seconds taken by all reduce tasks=40023040
```

As expected Mappers took much less time with Stripes compared to pairs as expected but in reducers Pairs took bit longer. This could be due to all unpacking we have to do with stripes to calculate final count.

END STUDENT CODE HW35

OPTIONAL

QUESTIONS BELOW THIS LINE ARE OPTIONAL

HW3.6 Computing Relative Frequencies on 100K WikiPedia pages (93Meg)

Dataset description For this assignment you will explore a set of 100,000 Wikipedia documents:

https://www.dropbox.com/s/n5lfbnztclo93ej/wikitext_100k.txt?dl=0 (https://www.dropbox.com/s/n5lfbnztclo93ej/wikitext_100k.txt?dl=0) s3://cs9223/wikitext_100k.txt, or https://s3.amazonaws.com/cs9223/wikitext_100k.txt (https://s3.amazonaws.com/cs9223/wikitext_100k.txt) Each line in this file consists of the plain text extracted from a Wikipedia document.

Task Compute the relative frequencies of each word that occurs in the documents in wikitext_100k.txt and output the top 100 word pairs sorted by decreasing order of relative frequency.

Recall that the relative frequency (RF) of word B given word A is defined as follows:

$$f(B|A) = \text{Count}(A, B) / \text{Count}(A) = \text{Count}(A, B) / \sum_B'(\text{Count}(A, B'))$$

where count(A,B) is the number of times A and B co-occur within a window of two words (co-occurrence window size of two) in a document and count(A) the number of times A occurs with anything else. Intuitively, given a document collection, the relative frequency captures the proportion of time the word B appears in the same document as A. (See Section 3.3, in Data-Intensive Text Processing with MapReduce).

In the async lecture you learned different approaches to do this, and in this assignment, you will implement them:

- Write a mapreduce program which uses the Stripes approach and writes its output in a file named rfstripes.txt
- Write a mapreduce program which uses the Pairs approach and writes its output in a file named rfpairs.txt
- Compare the performance of the two approaches and output the relative performance to a file named rfcomp.txt. Compute the relative performance as follows: (running time for Pairs/ running time for Stripes). Also include an analysis comparing the communication costs for the two approaches. Instrument your mapper and reducers for counters where necessary to aid with your analysis.

NOTE: please limit your analysis to the top 100 word pairs sorted by decreasing order of relative frequency for each word (tokens with all alphabetical letters).

Please include markdown cell named rf.txt that describes the following:

the input/output format in each Hadoop task, i.e., the keys for the mappers and reducers the Hadoop cluster settings you used, i.e., number of mappers and reducers the running time for each approach: pairs and stripes

You can write your program using Python or MrJob (with Hadoop streaming) and you should run it on AWS. It is a good idea to develop and test your program on a local machine before deploying on AWS. Remember your notebook, needs to have all the commands you used to run each Mapreduce job (i.e., pairs and stripes) -- include the Hadoop streaming commands you used to run your jobs.

In addition the All the following files should be compressed in one ZIP file and submitted. The ZIP file should contain:

A. The result files: rfstripes.txt, rfpairs.txt, rfcomp.txt

Prior to working with Hadoop, the corpus should first be preprocessed as follows: perform tokenization (whitespace and all non-alphabetic characters) and stopword removal using standard tools from the Lucene search engine. All tokens should then be replaced with unique integers for a more efficient encoding.

== Preliminary information for the remaining HW problems==

Much of this homework beyond this point will focus on the Apriori algorithm for frequent itemset mining and the additional step for extracting association rules from these frequent itemsets. Please acquaint yourself with the background information (below) before approaching the remaining assignments.

=== Apriori background information ===

Some background material for the Apriori algorithm is located at:

HW3.7 Apriori Algorithm

What is the Apriori algorithm? Describe an example use in your domain of expertise and what kind of . Define confidence and lift.

NOTE: For the remaining homework use the online browsing behavior dataset located at (same dataset as used above):

<https://www.dropbox.com/s/zlfyiwa70pogg74/ProductPurchaseData.txt?dl=0>

Each line in this dataset represents a browsing session of a customer. On each line, each string of 8 characters represents the id of an item browsed during that session. The items are separated by spaces.

Here are the first few lines of the ProductPurchaseData FRO11987 ELE17451 ELE89019 SNA90258 GRO99222 GRO99222 GRO12298 FRO12685 ELE91550 SNA11465 ELE26917 ELE52966 FRO90334 SNA30755 ELE17451 FRO84225 SNA80192 ELE17451 GRO73461 DAI22896 SNA99873 FRO86643 ELE17451 ELE37798 FRO86643 GRO56989 ELE23393 SNA11465 ELE17451 SNA69641 FRO86643 FRO78087 SNA11465 GRO39357 ELE28573 ELE11375 DAI54444

HW3.8. Shopping Cart Analysis

Product Recommendations: The action or practice of selling additional products or services to existing customers is called cross-selling. Giving product recommendation is one of the examples of cross-selling that are frequently used by online retailers. One simple method to give product recommendations is to recommend products that are frequently browsed together by the customers.

Suppose we want to recommend new products to the customer based on the products they have already browsed on the online website. Write a program using the A-priori algorithm to find products which are frequently browsed together. Fix the support to $s = 100$ (i.e. product sets need to occur together at least 100 times to be considered frequent) and find itemsets of size 2 and 3.

Then extract association rules from these frequent items.

A rule is of the form:

$(\text{item1}, \text{item5}) \Rightarrow \text{item2}.$

List the top 10 discovered rules in decreasing order of confidence in the following format

$(\text{item1}, \text{item5}) \Rightarrow \text{item2}, \text{supportCount}, \text{support}, \text{confidence}$

HW3.8.1

Benchmark your results using the pyFIM implementation of the Apriori algorithm (Apriori - Association Rule Induction / Frequent Item Set Mining implemented by Christian Borgelt). You can download pyFIM from here:

<http://www.borgelt.net/pyfim.html> (<http://www.borgelt.net/pyfim.html>)

Comment on the results from both implementations (your Hadoop MapReduce of apriori versus pyFIM) in terms of results and execution times.

END OF HOMEWORK